

Detecting AI-Generated Video: A Vision-Language Dual-View Survey

Dylan Xinming Hou¹, Juntian Zhang², Xu Gu²,
Yichen Wu³, Nils Lukas¹, Gus Xia¹, Xiuying Chen¹, Yuhan Liu^{1†}

¹MBZUAI, ²Gaoling School of Artificial Intelligence, Renmin University of China,

³Harvard University

dыхhou@gmail.com, {nils.lukas, gus.xia, xiuying.chen, yuhan.liu}@mbzuai.ac.ae,
{zhangjuntian, guxu}@ruc.edu.cn, yiwu6@mh.harvard.edu

The evolving realism of AI-generated Videos (AIGC-V) is rapidly rendering traditional artifact-centric detection insufficient, necessitating a paradigm shift from low-level inspection to high-level semantic verification. This paper presents a comprehensive survey of AIGC-V detection, reframing the task as Factual Fidelity Verification, which asks whether the events, entities, and physical processes depicted in a video are consistent with real-world facts. To systematize this rapidly evolving field, we propose a *Vision-Language Dual-View* taxonomy that organizes existing methods into a hierarchical, four-layer landscape, spanning intrinsic cue analysis, spatiotemporal consistency modeling, cross-modal consistency reasoning, and language-guided world-level reasoning. This dual-view framing highlights a fundamental transition from artifact matching in traditional deepfake detection to evidence-based semantic verification enabled by vision-language models and agentic reasoning pipelines. Based on a systematic review of 221 works, we synthesize AIGC-V generation paradigms, survey the landscape of detection methods, and review evaluation metrics and benchmarks in line with proposed views. Finally, we discuss current challenges and identify promising directions toward robust, explainable, and trustworthy detection.

Date: April 14, 2026

🔗 Github: <https://github.com/dxhou/AI-Generated-Video-Detection>

🏠 Homepage: <https://aigcvdetection.github.io/>



1 Introduction

The rapid evolution of video generation models, exemplified by Sora 2 (OpenAI, 2024b), Veo 3 (Google DeepMind, 2025), and Seedance 2.0 (ByteDance Seed, 2026), is fundamentally reshaping the information landscape. Unlike early deepfakes dominated by localized face swapping (Wang et al., 2024b), modern **AI-Generated Content-Videos (AIGC-V)** have achieved cinematic fidelity with coherent narratives, increasingly blurring the boundary between synthesized fiction and captured reality. This technological leap destabilizes the foundational trust in video evidence, traditionally relied upon to verify "who, when, where, and what" (Ho et al., 2022a).

†Corresponding author: Yuhan Liu(yuhan.liu@mbzuai.ac.ae)




Figure 1 An example of the AIGC-V detection pipeline under our approach, illustrating AI-generated videos from traditional methods or text-to-video prompts, detection from the visual and language views, and outputs at different levels.

As generation paradigms shift from local manipulation to end-to-end synthesis, such as text-to-video, traditional detection methods (Wang et al., 2025d; Ma et al., 2025) face a critical bottleneck. Early detection systems primarily relied on low-level visual artifacts, such as blending boundaries. However, advanced diffusion models and transformers (Ho et al., 2022b; OpenAI, 2024a) can now produce visually high-fidelity videos. This paradigm shift necessitates a new detection landscape, shifting from perceptual inspection (checking for visual artifacts) to cognitive reasoning (checking for semantic and factual violations). The surge of Vision-Language Models (VLMs) (Zhang et al., 2025d,c; Wang et al., 2026) and agentic frameworks (Liu et al., 2024b, 2025g,f; Hou et al., 2024) offers a promising pathway, enabling detectors to perform world-level reasoning and verify cross-modal consistency. Although Fu et al. (2025) and Park et al. (2025) have begun to explore these directions, the existing literature remains fragmented. Meanwhile, surveys like Pei et al. (2024) and Liu et al. (2025c) are constrained to the early era of deepfakes or treat video merely as a sequence of images, failing to systematize the emerging class of methods that leverage language for semantic verification.

To address this gap, we propose a *Vision-Language Dual-View* taxonomy to organize the AIGC-V detection landscape, as illustrated in Figure 1. We reframe the problem as *Factual Fidelity Verification*, determining whether the video content aligns with real-world facts and physical laws, rather than simple binary classification. Our framework categorizes detection methods into a four-layer hierarchy, progressing from low-level perception to high-level cognition: (1) Intrinsic Cue Analysis, (2) Spatiotemporal Consistency, (3) Cross-Modal Consistency, and (4) Language-Guided World-Level Reasoning.

In this survey, we comprehensively review 221 works as of March 2026. We begin with AIGC-V paradigms (§2), then formulate AIGC-V detection from the perspective of *factual fidelity* and introduce our Vision-Language dual-view framing (§3). We systematically organize methods under this landscape (§4), revisit evaluation metrics (§5.1) from a dual-view, and review benchmarks in line with AIGC-V paradigms (§5.2). Finally, we identify the critical challenges (§6) in detecting increasingly sophisticated AIGC-V. By bridging the gap between traditional visual forensics and emerging multimodal reasoning, this paper aims to provide a structured roadmap for the next generation of explainable and trustworthy AIGC-V detection.



Figure 2 Overview of three AIGC-V paradigms defined in this survey, highlighting typical forms.

2 Paradigms of AI-Generated Video

Currently, diffusion-based and Transformer-based architectures are advancing rapidly in video generation, driving major progress in text-to-video and other AIGC-V pipelines (Ho et al., 2022b,a; Singer et al., 2022; OpenAI, 2024b; Google DeepMind, 2025; ByteDance Seed, 2026). In this survey, we group mainstream methodologies into three categories according to the underlying generation paradigm: local manipulation, audio-visual editing, and generative video synthesis, following representative systems in each regime (Brison et al., 2025; Cheng et al., 2022b; Ho et al., 2022b). Figure 2 provides an overview.

2.1 Local Manipulation

Methodologies categorized as *Local Manipulation* typically operate on authentic video sequences by modifying specific spatial regions or distinct semantic attributes (Brison et al., 2025; Heo and Woo, 2025). The resulting generation maintains a high degree of structural fidelity to the original footage. Another technical paradigm centers on video face swapping and facial element manipulation. Representative approaches typically utilize diffusion models (Wang et al., 2024b), 3D facial priors (Wang et al., 2025d), and controllable conditional encoding, integrating identity feature transfer with expression, pose, and illumination reconstruction from source videos into a unified generative framework. Given their ease of implementation and covert nature, these techniques represent a critical threat in real-world adversarial scenarios.

2.2 Audio-Visual Editing

Audio-Visual Editing uses speech as an explicit control signal to drive talking-head generation or to re-dub existing videos, requiring tight cross-modal alignment between audio and face dynamics while preserving identity and background (Cheng et al., 2022b). And recent diffusion-based or 3D-aware approaches (Ma et al., 2025; Hong et al., 2025) move toward unified conditional generation, where audio guidance and identity constraints are jointly modeled to produce lip-synced (Li et al., 2024a), visually consistent outputs. Due to their low operational barrier and strong perceptual plausibility, audio-visual edits can convincingly “bind” forged visual performances to plausible voice tracks, amplifying risks in impersonation and misinformation.

2.3 Generative Video Synthesis

Generative Video Synthesis targets end-to-end creation of complete video sequences from text or noise, fabricating both appearance and dynamics without authentic carriers. Diffusion-based generators extend image diffusion with temporal modeling for coherence, exemplified by Video Diffusion Models (Ho et al., 2022b) and cascaded high-fidelity pipelines such as Imagen Video (Ho et al., 2022a); Make-A-Video (Singer et al., 2022) further reduces reliance on paired text-video data by leveraging text-image priors and unpaired videos. Recent architectures further optimize representations: Show-1 (Zhang et al., 2025b) hybridizes pixel-based generation with latent refinement for efficiency, while Grid Diffusion (Lee et al., 2024) unifies spatiotemporal dimensions into a single 2D grid to simplify modeling. Transformer-style synthesis tokenizes videos and learns long-horizon structure via autoregressive or masked modeling, while DiT further highlights the scalability of Transformer backbones for diffusion-based generation (Peebles and Xie, 2023). At the industrial frontier, large-scale text-to-video systems such as Sora 2 (OpenAI, 2024b), Seedance 2.0 (ByteDance Seed, 2026), and Veo 3 (Google DeepMind, 2025) demonstrate strong open-domain capability and raise urgent needs for provenance and verification.

Where the Evidence Lives

- ▶ **LMV keeps a real carrier**, so the strongest clues are usually residual and spatially localized rather than scene-wide.
- ▶ **AVE is constrained by coupling**, making synchrony, speaker consistency, and speech-conditioned facial motion more decisive than generic image artifacts.
- ▶ **GVS removes the carrier altogether**, shifting detection away from local edit traces toward long-range coherence, factual verification, and provenance.

3 AI-Generated Video Detection

3.1 Task Scope

Visually high fidelity AIGC-V with richer narratives make the reliability of methods that output only a “real/fake” probability increasingly low in security and privacy scenarios (Wang et al., 2024a; Kaur et al., 2024), and the detection task should shift from merely answering the traditional question of “whether the video is AI generated” to **fact-level** judgment of whether the video content “objectively happened in the real world.”

From a fact-level perspective, video content can be abstracted as a series of propositions of the form “at a certain time and place, which entities exist, and what happens to these entities.” We define **Factual Fidelity** as “the fact-level propositions implied by the video and its metadata are consistent with the real world,” and treat *Factual Fidelity Verification* as the **fundamental objective of AIGC-V detection**.

Furthermore, different evidence pathways provided by detection methods are integrated to indicate at which levels the video violates factual fidelity (Zhang et al., 2024b; Yu et al., 2025; Shen et al., 2025) and form an interpretable and trustworthy evidence space.

Definition 3.1 (Factual fidelity verification). Let $x \in \mathcal{X}$ be a video sample together with its




Figure 3 Dual-view, four-layer framework for AIGC-V detection. The visual view, spanning Layers 1-2, models intrinsic cues and spatiotemporal consistency, whereas the language view, spanning Layers 3-4, addresses cross-modal consistency and world-level reasoning, with each layer further decomposed into fine-grained subcategories. Boundary note: Layer 1 uses frame-level distributional evidence; Layer 2 models frame-sequence relations; Layer 3 performs within-video multimodal verification; Layer 4 requires external knowledge for verification.

metadata, and let $\Pi(x) \subseteq \mathcal{P}$ denote the set of fact-level propositions implied by x . Define the ideal factual-violation label as $y(x) = \mathbf{1}[\Pi(x) \not\subseteq \mathcal{P}_{\text{world}}]$, i.e., $y(x) = 1$ if and only if x violates factual fidelity.

Definition 3.2 (AIGC-V detection). An AIGC-V detector is a mapping $D : \mathcal{X} \rightarrow [0, 1] \times \mathcal{E}$. For a video sample $x \in \mathcal{X}$, $D(x)$ contains a suspicion score $s(x) \in [0, 1]$ and evidence $e(x) \in \mathcal{E}$ supporting its judgment.

We require $e(x)$ to be *trustworthy* and *grounded*. Concretely, we recommend representing $e(x)$ as a finite set of evidence items $e(x) = \{\epsilon_k\}_{k=1}^K$, where each $\epsilon_k = (g_k, c_k, \pi_k, r_k)$ contains: (a) a grounding g_k , such as temporal segment and optional spatial track, (b) a confidence $\pi_k \in [0, 1]$, (c) a possible human-readable claim-check statement c_k describing what sub-claim is being verified and what inconsistency is found, and (d) a provenance record r_k such as tool outputs or retrieved facts enabling traceability.

3.2 Dual-view Four-Layer Taxonomy

The methodological perspectives of AIGC-V detection can be summarized as two complementary scientific views: visual and language. Out of this observation, as illustrated in Figure 3, we propose a **Vision-Language Dual-View** approach and organize existing methods into a four-layer methodological landscape from low-level perception to high-level cognition in §4.

Definition 3.3 (Multimodal video sample). *Let $\mathcal{I} \subseteq \mathbb{R}^{H \times W \times C}$ denote the space of frames, and let $\mathcal{A}, \mathcal{S}, \mathcal{M}$ denote the spaces of audio, text, and metadata. A multimodal video sample is*

$$x = (V, A, S, M) \in \mathcal{X} := \mathcal{I}^T \times \mathcal{A} \times \mathcal{S} \times \mathcal{M},$$

where $V = (I_t)_{t=1}^T$ is the visual stream, and A, S, M may be empty when the corresponding modality is unavailable.

Visual View. The visual view focuses on the visual modality and emphasizes statistical differences between AIGC-V and real videos (Zhou et al., 2024; Gu et al., 2021). The detection pathway extends from low-level intrinsic cue analysis at the frame level in §4.1 to spatiotemporal consistency across frames in §4.2, forming the perceptual perspective for factual-fidelity verification.

Language View. In this survey, the language view denotes a *grounded verification pathway*, not an assumption that all audio cues are linguistic by default. Concretely, audio evidence enters this pathway in two forms. The first form is non-linguistic cross-modal perceptual evidence, which mainly checks within-video alignment without semantic parsing, such as lip–speech synchrony and onset and offset consistency (Chugh et al., 2020; Yang et al., 2021). The second form is language-grounded speech evidence: speech can be mapped via ASR or speech representations to phoneme, word, or semantic units, enabling consistency checks against identity and text content, such as voiceprint–identity coherence and speech–subtitle agreement (Shahzad et al., 2022; Bohacek and Farid, 2024; Cheng et al., 2022a; Zhang et al., 2025f). This design choice clarifies why audio can legitimately participate in the language view: it provides an interface from acoustic signals to language-level evidence and textual reasoning when needed. Building on these two cue types, methods progress from cross-modal consistency analysis (§4.3) to world-level reasoning at factual and knowledge levels when external evidence is required (§4.4) (Zhang et al., 2024b; Yu et al., 2025; Shen et al., 2025), thereby forming the cognitive perspective of factual-fidelity verification.

What AIGC-V Detection Must Establish

- ▶ **The target is factual fidelity, not only a real/fake score:** detection should judge whether the fact-level claims implied by a video and its metadata remain consistent with the real world.
- ▶ **A detector should return evidence, not only suspicion:** the output should couple a score with grounded, traceable evidence showing which claim is being checked, where the inconsistency lies, and why the decision is reliable.
- ▶ **Evidence can arrive through two complementary pathways:** the visual view tests intrinsic and temporal plausibility, while the language view moves from within-video cross-modal verification to explicit claim checking and world-level reasoning.




Figure 4 Landscape of representative AIGC-V detection methods aligned with the four-layer taxonomy in §4. Methods are grouped by the evolving capabilities of generative models that pose emerging threats and the corresponding advances in foundation models that enable detection. The year-wise shift toward language-view methods across the surveyed detection papers is quantified in Table 5.

4 Landscape of Detection Methods

As illustrated in Figure 4, we organize AIGC-V detection methods into four layers, and provide a structured overview in this section. Our taxonomy is **evidence-oriented rather than task-exclusive**: layer assignment follows the dominant evidence pathway used for decision-making, not a hard partition of tasks. To make layer assignment operational, we use boundary criteria tied to evidence dependency: Layer 1 captures frame-level distributional cues that can be scored per frame and pooled, whereas Layer 2 requires explicit frame-sequence modeling of spatiotemporal relations. Layer 3 then moves from visual-only coherence to within-video multimodal verification, whereas Layer 4 further verifies claims against external facts, commonsense, or physics, even when within-video modalities appear mutually consistent. When a method combines cues from multiple layers, we map each module to its primary evidence type rather than forcing a single-task label. For each layer below, we further summarize the assumption-signal-failure-mode linkage to clarify why each pathway works, when it breaks, and how recent methods improve robustness under the factual-fidelity framing in Section 3. To make this trend concrete, Table 5 summarizes the surveyed detection papers in Section 4 using complete publication years through 2025.

4.1 Layer 1: Intrinsic Cues Analysis

Layer 1 asks whether low-level visual signals obey the statistical regularities of real videos, whether they match acquisition and post-processing effects (Le and Woo, 2023; Wang et al., 2022b), and whether AI generation or editing leaves intrinsic cues such as style patterns, model artifacts (Chen et al., 2021), or deviant physiological signals (Ciftci et al., 2020b; Qi et al., 2020; Ciftci et al., 2020a). Methods operate from the visual view by modeling, extracting, and amplifying these low-level signals (Vahdati et al., 2024). The dominant evidence unit remains frame-level and distributional: a single frame can already yield an independent signal, while multiple frames mainly aggregate per-frame evidence. Once the decision depends on explicit inter-frame relations, the evidence moves to Layer 2. Table 1 summarizes representative methods in this layer.

Pixel and Geometric Artifacts. This line of work targets intrinsic traces introduced by AIGC-V synthesis and post-processing, including frequency fingerprints, texture residues, and geometric inconsistencies. FreqBlender amplifies synthesis-related spectral cues via spectral blending and contrastive learning (Zhou et al., 2024). Such frame-wise frequency fingerprints remain Layer 1 evidence, whereas descriptors that depend on temporal-frequency behavior belong to Layer 2. Localized designs reduce reliance on global spectra by focusing on region-level residues, such as mask-guided localization in MagDR and texture-consistency learning in HCIL (Chen et al., 2021; Gu et al., 2022b). Geometric cues complement spectral and texture cues by testing structural plausibility, as instantiated by LRNet, RAM, and head-pose inconsistency tests (Sun et al., 2021; Tian et al., 2024; Yang et al., 2019). Following quality-aware training results (Le and Woo, 2023), these cues still degrade under compression and adaptive filtering in low-quality pipelines.

Physiological Features. Physiological cues treat faces as carriers of biological rhythms and micro-dynamics that are difficult to synthesize consistently. Because many practical AIGC-V scenarios remain human-centric, especially talking heads and avatars, we keep physiological evidence as a dedicated subclass. Recent follow-up is still limited: beyond classic rPPG pipelines, only a few works revisit local-global physiological interaction patterns (Wu et al., 2024; D’Amelio et al., 2023). Heart-related signals are extracted either by mapping residual representations to heart-rate domains (Ciftci et al., 2020b) or by learning spatiotemporal attention over rPPG-related patterns (Qi et al., 2020). Higher-order heart-rate variability and cross-region coupling improve robustness under noise (Hernandez-Ortega et al., 2020; Fernandes et al., 2019; Lee et al., 2021; Wu et al., 2024; Stefanov et al., 2022; Mao and Yang, 2021; Ciftci et al., 2020a). Complementary evidence includes blink irregularities and identity-specific muscle dynamics (Li et al., 2018; Cozzolino et al., 2021). These cues require stable facial visibility and sufficient temporal support, and remain sensitive to occlusion and aggressive re-encoding.

Representation Distribution Discrepancy. This line of work targets cross-generator transfer by modeling distribution gaps between authentic and synthetic videos rather than committing to a fixed artifact family. Latent and style-space perturbations diversify synthetic styles and reduce forgery specificity (Yan et al., 2024), while latent-flow modeling

captures abnormal style trajectories (Choi et al., 2024). Curriculum learning and synthetic Blend-Fake data expose models to controlled difficulty and reduce shortcut learning (Lin et al., 2024b; Cheng et al., 2024). Recent follow-up stays close to low-level evidence: DFA combines foundation-model global priors with explicit local anomaly analysis to improve generalization under evolving forgery patterns (Liao et al., 2026). Discrepancy objectives separate overlapping real and fake manifolds with bounded contrastive learning and explicitly modeled appearance shifts (Larue et al., 2023; Liu et al., 2024a). At inference time, one-shot adaptation (Chen et al., 2022), quality-aware training (Le and Woo, 2023), and locality-aware reconstruction (Du et al., 2020) improve robustness under shift while also surfacing region-level evidence for review.

Takeaways for Layer 1

- **Why It Works:** Low-level statistics still preserve residual capture and synthesis traces, so frame-wise forensic evidence can separate authentic from generated content.
- **Paradigm Fit:** Strongest in LMV; in AVE it is usually auxiliary unless the visual stream is also manipulated; in GVS, generator fingerprints help but transfer weakens across models, pipelines, and quality conditions.
- **Failure Modes:** Compression, transcoding, and domain shift attenuate these traces, and physiology-based cues are even more brittle because they require clear faces, stable signal quality, and favorable capture conditions.

Method	Cue	Input	Mechanism	Output	Date
A. Pixel and geometric artifacts					
Inconsistent Head Poses (Yang et al., 2019)	3D head-pose inconsistency	V	Head-pose estimation + geometry checks	Score	05/2019
MagDR (Chen et al., 2021)	Localized artifacts around manipulated regions	F+V	Mask-guided localization + reconstruction loss	Score+loc.	06/2021
HCIL (Gu et al., 2022b)	Region-level inconsistency	V	Hierarchical contrastive learning	Score	10/2022
NoiseDF (Wang and Chow, 2023)	Forensic noise traces (face vs. bg)	F	Denoiser-based noise extraction + fusion	Score	06/2023
FreqBlender (Zhou et al., 2024)	Frequency-domain fingerprints	F	Spectral blending augmentation for frequency cues	Score	12/2024
B. Physiological features					
In Ictu Oculi (Li et al., 2018)	Eye-blink irregularities	V	Blink detection + temporal pattern modeling	Score	12/2018
FakeCatcher (Ciftci et al., 2020a)	PPG-like biological signal maps	V	Biological signal maps + detector	Score	07/2020
Hearts (Ciftci et al., 2020b)	Heart-related signals in residuals	V	Residual maps → rPPG features	Score	09/2020
DeepRhythm (Qi et al., 2020)	Visual heartbeat rhythms (rPPG)	V	Spatiotemporal attention over rPPG patterns	Score	10/2020
DeepFakesON-Phys (Hernandez-Ortega et al., 2020)	Heart-rate estimation inconsistency	V	rPPG heart-rate estimation + anomaly cues	Score	10/2020
Local rPPG Interaction (Wu et al., 2024)	Cross-region rPPG coupling	V	Local attention + long-range rPPG interaction	Score	02/2024
C. Distribution discrepancy and robustness					
OST (Chen et al., 2022)	Domain shift and compression sensitivity	F	One-shot test-time adaptation	Score	12/2022
SeeABLE (Larue et al., 2023)	Soft discrepancies under shifts	F	Real-only bounded contrastive learning	Score	10/2023
QAD (Le and Woo, 2023)	Compression-robust representations	F	Quality-aware regularization	Score	10/2023
LSDA (Yan et al., 2024)	Cross-generator transfer	V	Latent-space augmentation for transfer	Score	06/2024
Style Latent Flows (Choi et al., 2024)	Abnormal style-latent trajectories	V	Style-flow modeling + contrastive loss	Score	06/2024
Fake It Till You Make It (Lin et al., 2024b)	Curriculum-based forgery augmentation	V	Curriculum forgery augmentation	Score	10/2024
DFA (Liao et al., 2026)	Global-local intrinsic anomalies	F+V	CLIP adapter + local anomaly stream	Score	03/2026

Table 1 Layer 1: Intrinsic cues analysis. Representative methods grouped by cue family. *Cue* = evidence source; *Mechanism* = modeling design; Input: F=frames, V=video.

4.2 Layer 2: Spatiotemporal Consistency

Layer 2 asks whether spatiotemporal image flow satisfies the motion constraints of real videos. Real capture is constrained by continuous camera trajectories and physical scenes (Zhang et al., 2025e; Internò et al., 2025), so adjacent frames exhibit continuous, predictable, and physically feasible changes, whereas AIGC-V can show long-range inconsistencies such as object or background distortion (Gu et al., 2022c; Zhang et al., 2024a) and sudden local blurring (Gu et al., 2022a). Spatiotemporal inconsistency learning (Gu et al., 2021) and motion-based detection (Ma et al., 2024b; Yin et al., 2023; Haliassos et al., 2021) therefore exploit frame differences and flow residuals (Chen et al., 2025b). This makes Layer 2 distinct from Layer 1 by requiring explicit sequence modeling of inter-frame relations rather than frame-wise cue pooling. At the same time, Layer 2 remains a visual-modality layer: once verification requires semantic checks across audio, text, and visual events, the evidence pathway transitions to Layer 3. Table 2 summarizes representative methods in this layer.

Temporal and Motion Inconsistencies. Layer 2 makes explicit how temporal coherence is operationalized as a modeling prior. Instead of classifying frames independently, spatiotemporal methods model how facial appearance and geometry evolve across time and treat incoherence as evidence for forgery (Gu et al., 2021, 2022c). Clip encoders treat short sequences as space-time volumes and learn joint patterns with 3D CNNs or attention (Gu et al., 2021, 2022c; Hu et al., 2021; Zhang et al., 2024a; Chen et al., 2020; Zhang et al., 2021a). By jointly observing multiple frames, they reveal abnormal changes that are weak or ambiguous at the single-frame level. Region-aware variants highlight regions such as the eyes and mouth where features change abnormally, and auxiliary objectives can predict inconsistency maps to force localized temporal evidence (Gu et al., 2022c; Hu et al., 2021; Zhang et al., 2024a). Temporal Dropout and rhythm perturbation further reduce overfitting to specific frame rates or speaking rhythms (Chen et al., 2020; Zhang et al., 2021a). To reuse 2D backbones, some methods reshape frame sequences into grid-like layouts to enable efficient 2D processing and leverage image pre-training (Xu et al., 2024, 2023). Hybrid pipelines aggregate per-frame features with recurrent branches to fuse static appearance and dynamic motion cues at the decision stage (Masi et al., 2020). Application-oriented follow-ups continue to revisit CNN or LSTM aggregation as a lightweight temporal baseline for video deepfake detection (Bahadure and Mane, 2026). Motion-centric designs construct intermediate signals such as frame differences or optical flow to emphasize temporal transitions (Ma et al., 2024b; Yin et al., 2023). Flow-residual cues test whether local facial motion is compatible with global head motion, remaining informative under blur and compression (Chen et al., 2025b; Wang et al., 2023). Overall, spatiotemporal inconsistency methods frame deepfake detection as coherence analysis rather than static appearance classification.

Human Behavioral and Interaction Dynamics. Behavioral cues move from artifact matching to higher-level human realism, asking whether expressions, gaze, and identity evolve plausibly over time. Action-unit guided and appearance-behavior joint representations model muscle activations and expression trajectories to expose abnormal combinations

or unnatural dynamics (Anand et al., 2025; Agarwal et al., 2020). Gaze-based cues probe attention and interaction patterns that are difficult to reproduce consistently in face swapping or reenactment (Kohler et al., 2025). Identity-driven methods test whether temporal dynamics remain compatible with the claimed person under varying pose, illumination, and compression (Dong et al., 2020). Geometry-consistency checks enforce compatibility between 2D observations and 3D face and head motion, exposing pose- or shape-inconsistent sequences (Sun et al., 2021; Yang et al., 2019; Tursman et al., 2020). These semantics- and geometry-aware cues can be strong against visually polished fakes but are sensitive to natural behavioral diversity and recording conditions.

Physical and Frequency Artifacts. Transform-domain designs compute descriptors where non-physical dynamics become separable from content variation, yielding traces that can be more stable than raw RGB. Compared with purely spatial or spatiotemporal backbones trained on RGB, these descriptors can be more robust to appearance changes and sometimes more transferable across generators (Kim et al., 2025). Temporal frequency responses expose abnormal periodic components and energy distributions along time, helping separate authentic motion spectra from synthetic artifacts (Nie et al., 2024; Kim et al., 2025). Early-2026 work pushes this line toward higher-fidelity synthetic videos: MPF-Net combines static off-manifold deviation scoring with micro-temporal fluctuation analysis, arguing that visually polished clips can still reveal structured instability around the real-video manifold (He et al., 2026). Second-order statistics capture covariance mismatches and subtle distribution shifts in spatiotemporal features that survive mild post-processing (Zheng et al., 2025; Padhi et al., 2025). Residual- and difference-based features amplify deviations from real-video manifolds by emphasizing temporal changes rather than absolute appearance (Xu et al., 2025d; Stamnas and Sanchez, 2025). Physics-inspired priors provide complementary constraints on motion feasibility and perceptual plausibility, linking detection to camera trajectories and physical scenes (Zhang et al., 2025e; Internò et al., 2025). Two representative detectors in this subclass are NSG-VD (Zhang et al., 2025e) and ReStraV (Internò et al., 2025). NSG-VD injects explicit physics constraints into spatiotemporal modeling by matching normalized spatiotemporal gradients and distribution statistics between real and generated videos, while ReStraV scores how far latent motion trajectories deviate from straight, natural perceptual manifolds.

Generalization and Robustness. Beyond cue design, robustness strategies aim to maintain performance under unseen generators and post-processing by changing objectives, architectures, and training protocols (Nguyen et al., 2025; Yan et al., 2025; Trinh et al., 2021; Coccomini et al., 2022; Tong et al., 2025). Vulnerability-aware objectives and plug-and-play adaptation modules steer backbones toward transferable spatiotemporal evidence (Nguyen et al., 2025; Yan et al., 2025). Unified or cross-model detectors reuse pretrained backbones and expand the hypothesis space to handle diverse manipulations and media types (Trinh et al., 2021; Veeramachaneni et al., 2025; Coccomini et al., 2022; Liu et al., 2025a; Tariq et al., 2021). Forensic augmentation and quality-centric curricula expand coverage of perturbations and hard cases (Corvi et al., 2025; Song et al., 2024a). Self-supervised video forensics learns normal dynamics via synchronization or anomaly objectives and then treats deviations as manipulations (Yang et al., 2025; Feng et al., 2023). Multiple-instance

formulations aggregate clip-level evidence to reduce reliance on exhaustive forged labels and handle mixed-quality videos (Li et al., 2020a).

Takeaways for Layer 2

- **Why It Works:** Real capture follows continuous trajectories and feasible motion, whereas generated video still accumulates subtle long-horizon temporal inconsistencies.
- **Paradigm Fit:** Especially useful in GVS and in LMV when edits disturb motion across frames; in AVE, it matters mainly when the visual stream itself is synthesized, such as lip reenactment.
- **Failure Modes:** Short clips, resampling, and compression dilute the signal, and detectors may still collapse to frame-level shortcuts instead of genuine temporal reasoning.

Method	Cue	Mechanism	Output	Date
A. Temporal and motion inconsistencies				
FS-Spotter (Chen et al., 2020)	Spatiotemporal swap artifacts	Fuse spatial+temporal cues	Score	07/2020
Dynamic Prototypes (Trinh et al., 2021)	Transferable spatiotemporal evidence	Dynamic prototypes + predictive learning	Score	01/2021
Temporal Dropout 3DCNN (Zhang et al., 2021a)	Frame-rate dependent temporal artifacts	3D CNN + temporal dropout	Score	08/2021
STIL (Gu et al., 2021)	Short-clip spatiotemporal inconsistency	Clip encoder for ST inconsistency	Score	10/2021
EfficientNet-ViT Ensemble (Coccomini et al., 2022)	Robust temporal features	CNN and ViT ensemble	Score	05/2022
Region-Aware Temporal Inconsistency (Gu et al., 2022c)	Region-wise dynamics anomalies	Region-aware learning + maps	Score+loc.	07/2022
TALL (Xu et al., 2023)	Cross-frame inconsistency	Thumbnail layout + 2D backbone	Score	10/2023
Natural Consistency Representation (Zhang et al., 2024a)	Temporal naturalness deviations	SSL consistency + fine-tune	Score	11/2024
Vulnerability-Aware Learning (Nguyen et al., 2025)	Cross-generator vulnerability patterns	Vulnerability-aware transfer	Score	01/2025
GC-ConsFlow (Chen et al., 2025b)	Optical-flow residual inconsistency	Flow residuals + global context	Score	06/2025
Plug-and-Play Adapters (Yan et al., 2025)	Generalization under generator shift	Video blending + ST adapters	Score	06/2025
B. Physical and frequency artifacts				
DIP (Nie et al., 2024)	Temporal frequency response anomalies	Transform-domain temporal spectra	Score	12/2024
ReStraV (Perceptual Straightening) (Internò et al., 2025)	Non-physical motion and trajectory irregularities	Perceptual straightening	Score	07/2025
Beyond RGB (Kim et al., 2025)	Stable descriptors beyond raw RGB	Transformed descriptors	Score	07/2025
Physics-Driven ST Modeling (Zhang et al., 2025e)	Physics constraints on spatiotemporal flow	Physics-driven ST modeling	Score	10/2025
MPF-Net (He et al., 2026)	Off-manifold residuals + micro-temporal fluctuations	Hierarchical manifold deviation + temporal filtering	Score	01/2026
C. Human behavioral and interaction dynamics				
Emotions Don’t Lie (Mittal et al., 2020)	Affective audio-visual cue coherence	Siamese audio-visual model + triplet loss	Score	03/2020
Appearance and Behavior (Agarwal et al., 2020)	Appearance and behavior realism over time	Behavior-aware temporal features	Score	12/2020
Identity-Driven Detection (Dong et al., 2020)	Identity dynamics consistency	Track identity dynamics	Score	12/2020
Gaze Tracking (Kohler et al., 2025)	Gaze dynamics	Gaze features for dyadic calls	Score	09/2025

Table 2 Layer 2: Spatiotemporal consistency. Representative methods grouped by (A) temporal and motion inconsistencies, (B) physical and frequency artifacts, and (C) human behavioral and interaction dynamics in §4.2. *Cue* = evidence source; *Mechanism* = modeling design. All methods take video as input.

4.3 Layer 3: Cross-Modal Consistency

Layer 3 focuses on whether the modalities in a video are well aligned in describing the same thing. Real videos often contain audio, text, and visuals that are highly aligned across modalities, whereas AIGC-V may present mismatches between lip movements and speech or between identity and voiceprint (Chugh et al., 2020; Yang et al., 2021; Zhang et al., 2021b; Cheng et al., 2022a). Mismatches can also appear between image content and associated text, motivating text-video consistency reasoning (Zhang et al., 2025f; Liu et al., 2025e; Wang et al., 2025a). Symbolic and semantic lip-sync analyses further enlarge this consistency space (Shahzad et al., 2022; Bohacek and Farid, 2024; Li et al., 2024b; Koutlis and Papadopoulos, 2025). Compared with Layer 2, the key shift is from visual-only sequence coherence to within-video multimodal verification; compared with Layer 4, evidence in Layer 3 is still primarily within-video and does not require external world knowledge as a mandatory source.

To clarify the role of audio in this layer, our taxonomy does not treat all audio cues as linguistic by default. Instead, we distinguish two cue types that naturally coexist in Layer 3. The first type is non-linguistic cross-modal perceptual cues, which test alignment without semantic parsing, such as lip-speech synchrony and onset and offset consistency. The second type is language-grounded speech cues, where speech is mapped via ASR or speech representations into phoneme, word, or semantic units, enabling explicit verification against identity and text content, such as voiceprint-identity coherence and speech-subtitle agreement, and providing a language interface for further reasoning when external knowledge is needed. Table 3 summarizes representative methods in this layer.

Audio-Visual Consistency Detection. Audio-visual consistency methods decompose evidence into synchrony, symbolic alignment, and identity coherence. Synchrony cues compare speech dynamics with mouth motion and learn cross-modal embeddings that penalize misalignment (Chugh et al., 2020; Yang et al., 2021). Cues extend beyond the mouth to additional regions whose micro-motions correlate with speech production and head dynamics (Agarwal and Farid, 2021). Symbolic alignment improves interpretability by mapping audio and visual streams to token or phoneme sequences and comparing them at the sequence level (Shahzad et al., 2022; Bohacek and Farid, 2024; Li et al., 2024b; Koutlis and Papadopoulos, 2025). Identity coherence shifts the question from “what is said” to “who is speaking” by testing whether voice and face embeddings belong to the same person, including person-of-interest settings trained with authentic talking-head data only (Zhang et al., 2021b; Cheng et al., 2022a; Cozzolino et al., 2023; Muppalla et al., 2023). ConLLM further couples contrastive multimodal alignment with LLM-based reasoning to reduce modality fragmentation and expose fine-grained inconsistencies that are difficult to separate with shallow fusion alone (Kashyap et al., 2026). Generator-side probes such as X-AVDT further exploit inversion-accessed cross-attention signals to expose hidden speech-motion correspondence errors, while AV-LMMDetect shows that large multimodal models can also be directly fine-tuned as within-video audio-visual verifiers without requiring external world knowledge (Kim et al., 2026; Cao et al., 2026).

Text-Video Semantic Consistency Reasoning. Text-video semantic consistency reasoning reframes detection as claim verification: given a video and its associated text, such as titles, captions, or OCR and ASR transcripts, the model tests whether the textual statements are supported by visual evidence, with semantic contradictions serving as a key signal. CA-FVD supervises cross-modal consistency by prompting a VLM to obtain pairwise modality-consistency pseudo labels across visual, text, and audio streams, then training with cosine-based consistency losses; co-attention fusion and collaborative diagnosis aggregate modality-level confidence into a final score (Zhang et al., 2025f). CSCL moves from global scoring to grounded evidence by building patch-token consistency matrices and using cascaded decoders to separately model within-modality context and cross-modality semantics; forgery-aware aggregation reduces confusing but locally consistent content and enables localization (Liu et al., 2025e). To handle event shifts, T³SVFND introduces test-time training with a reconstruction objective conditioned on multimodal context, suggesting event-robust text-video reasoning for detection settings (Wang et al., 2025a).

Robust Learning and Temporal Localization. In realistic settings, cross-modal inconsistencies are often temporally sparse, so localization is needed not only for final scoring but also for review and diagnosis. Robust methods therefore need to remain stable under compression, language shift, and partial-stream corruption while still marking when and where the mismatch emerges. This couples robustness with interpretability: a useful detector should preserve discrimination under noise while exposing the specific temporal span that carries the cross-modal failure rather than diffusing evidence over the whole clip. Two-stream fusion and attention-based alignment provide a common backbone for joint modeling and interpretation (Zhou and Lim, 2021; Wang et al., 2022a). Transferable representations learned from large-scale authentic data improve stability under compression and language shifts (Liang et al., 2025; Feng et al., 2023). Weak supervision with video-level labels and bias-aware objectives enables fine-grained temporal localization without dense annotations (Xu et al., 2025c; Astrid et al., 2025). Audio shift prediction provides an additional self-supervised proxy for sparse cross-modal disturbance and improves temporal localization in multimodal deepfake settings (Anshul et al., 2026). Challenge-driven systems such as Divide and Conquer (Li et al., 2026) also show the value of decomposing detection into audio-side and visual-side localization before late fusion, which is particularly useful when manipulations are sparse or concentrated in only one stream.

Takeaways for Layer 3

- ▶ **Why It Works:** Authentic videos obey strong within-video coupling across modalities, so mismatches in timing, identity, or semantics are hard to hide consistently.
- ▶ **Paradigm Fit:** Natural fit for AVE; useful in LMV when edits change speaker identity or event semantics and usable audio or text is present; in GVS, it is strongest for captioned or prompt-conditioned clips.
- ▶ **Failure Modes:** Its value drops when audio or text is missing, noisy, or weakly informative, and attackers can partially restore apparent alignment through redubbing, retiming, or other resynchronization.

Detecting AI-Generated Video: A Vision-Language Dual-View Survey

Method	Cue	Input	Mechanism	Output	Date
A. Audio-visual consistency detection					
Not Made for Each Other (Chugh et al., 2020)	Audio-visual dissonance	Speech+Face	Cross-modal mismatch + localization	Score+loc.	10/2020
Dynamic Lip Movement (Yang et al., 2021)	Lip motion vs. speech mismatch	Speech+Lip	Lip-motion dynamics analysis	Score	12/2020
Aural-Oral Dynamics (Agarwal and Farid, 2021)	Aural-oral dynamics coherence	Speech+Lip	Audio+mouth dynamics features	Score	06/2021
Joint Audio-Visual Detection (Zhou and Lim, 2021)	Audio-visual synchrony and semantics	Speech+Face	Joint synchrony+semantic modeling	Score	10/2021
Audio-Visual Attention (Wang et al., 2022a)	Cross-modal alignment	Speech+Face	Cross-attention fusion	Score	03/2022
Lip Sync Matters (Shahzad et al., 2022)	Symbolic lip-sync mismatch	Speech+Lip	Token and phoneme sequence comparison	Score	11/2022
Voice-Face Homogeneity (Cheng et al., 2022a)	Voice-face identity coherence	Voice+Face-ID	Compare voice and face-ID embeddings	Score	11/2023
Lost in Translation (Bohacek and Farid, 2024)	Language-aware audio-visual mismatch	Speech+Lip	Language-conditioned lip-sync	Score	06/2024
AVFF (Oorloff et al., 2024)	Intrinsic audio-visual correspondences	Speech+Face	Real-only SSL pretrain + classifier	Score	06/2024
Multi-task Audio-Visual Prompt Learning (Miao et al., 2025)	Fine-grained audio-visual consistency	Speech+Face	Prompting + matching loss + fusion	Score	04/2025
CAD (Du et al., 2025)	Semantic audio-visual misalignment + cues	Speech+Face	Alignment + distillation	Score	05/2025
PIA (Datta et al., 2025)	Phoneme-timing mismatch + ID dynamics	Speech+Lip	Phoneme timing + ID dynamics	Score	10/2025
KLASSify to Verify (Kukanov and Ng, 2025)	Robust audio-visual cues (unseen attacks)	Speech+Face	Audio SSL + graph attn + handcrafted	Score+loc.	10/2025
ConLLM (Kashyap et al., 2026)	Multimodal semantic inconsistency	Speech+Face	Contrastive alignment + LLM reasoning	Score	01/2026
X-AVDT (Kim et al., 2026)	Generator-internal audio-visual correspondence	Speech+Face	DDIM inversion + cross-attention probing	Score	03/2026
B. Text-video semantic consistency reasoning					
CA-FVD (Zhang et al., 2025f)	Video-text semantic mismatch	Video+Text	VLM pseudo-labels + consistency loss	Score	04/2025
CSCL (Liu et al., 2025e)	Text-image inconsistency grounding	Frame+Text	Cascaded consistency decoders	Score+loc.	06/2025
T³SVFND (Wang et al., 2025a)	Event-shifted semantics	Video+Text	Test-time MLM reconstruction	Score	07/2025
C. Robust learning and temporal localization					
Cross- and Within-Modality Regularization (Zou et al., 2024)	Modality separation under perturb.	Speech+Face	Audio-visual Transformer + cross and within regs	Score	04/2024
Audio-Visual Local Inconsistencies (Astrid et al., 2025)	Local temporal audio-visual inconsistency	Speech+Face	Local inconsistency + localization	Score+loc.	04/2025
Circumventing Shortcuts (Smeu et al., 2025)	Shortcut-robust audio-visual reps (silence)	Speech+Face	Real-only SSL audio-visual alignment	Score	06/2025
SpeechForensics (Liang et al., 2025)	Real-only audio-visual repr. learning	Speech+Lip	SSL pretrain on real + fine-tune	Score	08/2025
WMMT (Xu et al., 2025b)	Weakly-sup. temporal loc.	Speech+Face	Multitask + MoE + deviation loss	Score+loc.	08/2025
HOLA (Wu et al., 2025)	Hierarchical audio-visual aggregation	Speech+Face	SSL pretrain + gated aggregation	Score	10/2025
Weakly-Supervised Temporal Localization (Xu et al., 2025c)	Sparse multimodal deviations	Speech+Face	Weak supervision + deviation modeling	Score+loc.	10/2025
A-V Shift Prediction (Anshul et al., 2026)	Temporal audio-visual perturbation cues	Speech+Face	Audio shift prediction + localization	Score+loc.	01/2026
Divide and Conquer (Li et al., 2026)	Cross-modal fusion + temporal localization	Speech+Face	Unimodal detectors + score fusion	Score+loc.	02/2026

Table 3 Layer 3: Cross-modal consistency. Representative methods grouped by (A) audio-visual alignment, (B) text-video semantics, and (C) robust learning and localization.

4.4 Layer 4: Language-Guided World-Level Reasoning

Layer 4 elevates detection from internal consistency within the video to consistency with world-level rules and knowledge. The research question shifts to whether the video content can truly exist in the real world and remain reasonable in semantic and factual dimensions. Real videos should be consistent with real-world facts, physical rules, domain knowledge, and basic common sense, whereas AIGC-V often fails to align fully with these constraints, creating the detection space exploited by this layer (Motamed et al., 2025b). Language serves as a natural interface for human-readable explanations, while baseline VLM prompting in zero-shot or few-shot settings provides a minimal route for detecting inconsistencies and producing textual rationales (Jia et al., 2024; Shahzad et al., 2025). Compared with Layer 3, the detector is no longer asked only whether modalities agree, but whether the implied claim remains tenable once external facts, commonsense, or physics are brought into the loop. Table 4 summarizes representative methods in this layer.

Language-Calibrated Forgery Representations. Language-calibrated designs inject prompts or textual priors into multimodal representations to reshape forgery-sensitive features without rebuilding the backbone. CPML aligns physiological pulse (rPPG) signals and facial landmark dynamics with prompts and enforces cross-quality and cross-modal consistency to stabilize physiological evidence under compression (Lai et al., 2024; D’Amelio et al., 2023). RepDFD reprograms pretrained vision-language models by freezing the backbone and learning input-side perturbations and adaptive prompts, while knowledge-guided variants build textual prototypes and uncertainty modeling for improved transfer (Lin et al., 2025a; Yu et al., 2025; Shen et al., 2025). Adapter-style reprogramming provides a lightweight alternative when full fine-tuning is infeasible (Shao et al., 2025). These methods are useful when external retrieval is unavailable but language priors still need to bind visual cues to explicit semantic hypotheses and human-readable defect categories.

Tool-Augmented Agentic Pipelines. Tool-augmented pipelines cast AIGC-V detection as explicit evidence gathering, where language guides what to inspect and which analyzer to call next. LAVID uses an observe-tool-integrate loop: it forms a coarse hypothesis, calls external tools, updates prompts based on intermediate outputs, and fuses multi-step evidence (Liu et al., 2025d). FakeHunter adds memory-anchored observe-think-act routines, including retrieval of semantic anchors and reuse of intermediate results across steps (Chen et al., 2025a). DAVID-XR1 and related pipelines emphasize semantic anchors and modular analyzers to support interpretable evidence composition (Gao et al., 2025). Multi-agent coordination further decomposes modality-specific evidence before fusion (Zaman et al., 2025). The attraction of this design is that the final verdict can be traced to retrieved facts, tool outputs, and intermediate judgments rather than a single opaque forward pass.

Post-Training for Explainable World-Level Reasoning. Post-training internalizes evidence selection and reasoning style so that inference depends less on external controllers. Within Layer 4, reinforcement learning and preference alignment encode how evidence should be gathered and how explanations should be structured, as explored by VidGuard-

R1, Veritas, and BusterX++ (Park et al., 2025; Tan et al., 2025; Wen et al., 2025b). DeeptraceReward learns complementary reward supervision aligned to human-perceived fakeness in AI-generated videos (Fu et al., 2025), and VideoVeritas further couples preference alignment with perception pretext reinforcement learning and fact-based reasoning for more balanced AIGC-V verification (Tan et al., 2026). Related efforts extend post-training to multi-stage reasoning and modular explanation under complex AIGC-V settings (Wen et al., 2025a; Sun et al., 2025; Li et al., 2025b; Chen et al., 2024b). The emphasis shifts from generic instruction-following to disciplined justification, including when to retrieve, when to localize uncertainty, and how to present compact but auditable evidence.

Takeaways for Layer 4

- **Why It Works:** It targets a harder bottleneck than within-video agreement: a clip may look self-consistent yet still fail once its implied claims are checked against facts, commonsense, causality, or physical rules.
- **Paradigm Fit:** Best for GVS and claim-centric LMV or AVE cases where the decision turns on external verification of actors, events, relations, or physical plausibility, especially when within-video modalities appear mutually consistent.
- **Failure Modes:** Verification depends on knowledge coverage, retrieval freshness, and correct claim decomposition, and VLMs can hallucinate, over-abstract, or miss subtle visual evidence while sounding confident.

Method	Base Model and System	Training	Output	Date
A. Prompts and adapters for representation calibration				
ChatGPT Detect (Jia et al., 2024; Shahzad et al., 2025)	ChatGPT	Prompt	Verdict+exp.	06/2024
CPML (Lai et al., 2024; D’Amelio et al., 2023)	rPPG+landmark encoder	Prompt-guided	Score	11/2024
DeepFake-Adapter (Shao et al., 2025)	ViT (frozen)	Adapter tuning	Frame label (agg.)	01/2025
RepDFD (Lin et al., 2025a)	Frozen VLM	Prompt tuning	Label	04/2025
AuthGuard (Shen et al., 2025)	Vision encoder + LLM	Cls+ITC; uncertainty	Verdict+exp.	06/2025
LVLMDFD (Yu et al., 2025)	LVLm + LLM	Detector+prompt learner	Verdict+loc.+exp.	07/2025
B. Tool-augmented agents for evidence gathering				
LAVID (Liu et al., 2025d)	LVLm agent + tools	Tool loop	Verdict+evidence	02/2025
DAVID-XR1 (Gao et al., 2025)	Video VLM	SFT	Loc.+defect reasoning	06/2025
FakeHunter (Chen et al., 2025a)	Vision+audio encoders + VLM agent	Agent+retrieval	Verdict+evidence	08/2025
DeepAgent (Zaman et al., 2025)	Multi-agent pipeline	Not stated	MM verdict	12/2025
C. Post-training, preferences and rewards				
X2-DFD (Chen et al., 2024b)	LLaVA + aux detectors	SFT (explainable data)	Verdict+exp.	10/2024
BusterX (Wen et al., 2025a)	Qwen2.5-VL	SFT → RL	Verdict+exp.	05/2025
BusterX++ (Wen et al., 2025b)	Qwen2.5-VL	RL → SFT → RL	Verdict+struct. exp.	07/2025
VERITAS (Tan et al., 2025)	InternVL3 and Qwen2.5-VL	SFT → MiPO → P-GRPO	Verdict+exp.	08/2025
DeeptraceReward (Fu et al., 2025)	VideoLLaMA3 and Qwen2.5-VL with reward modeling	Reward-model train	Reward	09/2025
VidGuard-R1 (Park et al., 2025)	Qwen2.5-VL	SFT → DPO → GRPO	Verdict+exp.	10/2025
EDVD-LLaMA (Sun et al., 2025)	Qwen2.5-7B + video encoder	SFT	Verdict+exp.	10/2025
Skyra and Skyra-RL (Li et al., 2025b)	Qwen2.5-VL	SFT → RL	Verdict+grounded reasoning	12/2025
VideoVeritas (Tan et al., 2026)	Video MLLM	Pref. align + PPRL	Verdict+reasoning	02/2026

Table 4 Layer 4: Language-guided world-level reasoning. Representative methods grouped by (A) prompts and adapters, (B) tool-augmented agents, and (C) post-training with preferences and rewards. *Base Model and System* reports backbone names when available.

4.5 Quantitative Summary of the Four-Layer Landscape

4.5.1 Year-wise Layer Distribution

Table 5 summarizes the surveyed detection papers discussed in Section 4 and tracks how the four-layer taxonomy evolves over time. We report both yearly share and cumulative share. The yearly ratio reflects release-time dynamics within each publication year, whereas the cumulative ratio shows how the full surveyed landscape shifts as methods accumulate. Because 2026 is still in progress, the table reports complete publication years through 2025; the cumulative denominator additionally includes the three surveyed papers before 2020. Two patterns stand out. First, the early landscape is still overwhelmingly visual-view: from 2020 to 2022, Layers 1 and 2 account for 30 of 37 surveyed methods, showing that the field was primarily organized around forensic traces and temporal coherence. Second, the transition toward the language view is gradual rather than monotonic. The yearly language-view share rises from 7.7% in 2020 to 26.7% and 22.2% in 2021–2022, reaches 40.0% in 2023, and then expands further in 2024 and 2025. The decisive break arrives in 2025, when 30 of 48 surveyed methods fall into the language view. The cumulative language-view share correspondingly grows from 6.3% to 40.3%, indicating that multimodal verification and world-level reasoning have moved from a marginal line of work to a major axis of recent AIGC-V detection. At the same time, the visual view remains substantial, with 18 methods still belonging to Layers 1 and 2 in 2025; low-level and temporal evidence are therefore not replaced, but increasingly complemented by higher-level verification.

Year	Layer Counts					View Counts		Language-View Share	
	L1	L2	L3	L4	Total	Visual	Language	Yearly	Cumulative
2020	5	7	1	0	13	12	1	7.7% 1 of 13	6.3% 1 of 16
2021	5	6	4	0	15	11	4	26.7% 4 of 15	16.1% 5 of 31
2022	4	3	2	0	9	7	2	22.2% 2 of 9	17.5% 7 of 40
2023	3	3	4	0	10	6	4	40.0% 4 of 10	22.0% 11 of 50
2024	9	5	4	3	21	14	7	33.3% 7 of 21	25.4% 18 of 71
2025	0	18	15	15	48	18	30	62.5% 30 of 48	40.3% 48 of 119
Δ	-5	+11	+14	+15	+35	+6	+29	+54.8	+34.0

Table 5 Year-wise quantitative summary for the surveyed detection papers discussed in Section 4. The visual view combines Layers 1 and 2, whereas the language view combines Layers 3 and 4. Share values are reported as a percentage with the corresponding count ratio. The table reports complete publication years through 2025; the cumulative denominator additionally includes the three surveyed papers before 2020. In the Δ row, differences are computed between 2020 and 2025, and the two share columns report percentage-point changes.

4.5.2 Layer-wise Performance Snapshot

To complement the taxonomy with a compact quantitative anchor, Table 6 collects reported AUC (%) for representative methods across the four layers. Because reporting protocols remain heterogeneous, each number is paired with its protocol label, either cross-dataset (CD) or in-domain (ID). The table should therefore be read as a protocol-aware snapshot rather than as a single cross-layer leaderboard. Three observations are useful when reading the table. First, visual-view methods remain highly competitive under standard CD evaluation: strong Layer 1 and Layer 2 systems still achieve high AUC on CDFv2 and retain reasonable transfer to DFDCP, which shows that low-level and temporal evidence continue to generalize well when benchmarks emphasize generator shift. Second, Layer 4 methods are already competitive under the same CD setup. RepDFD and LVLMDFD match or exceed lower-layer baselines on CDFv2 and DFDCP, suggesting that language-guided verification can be added without giving up benchmark strength. Third, the table also reveals an evaluation gap for higher layers: Layer 3 is represented here only by an ID result, and many recent multimodal or world-grounded systems still lack standardized CD reporting. This means the table should not be read as evidence that later layers uniformly outperform earlier ones. Rather, it shows that different layers answer different verification needs, while current benchmark practice is still better aligned with artifact and coherence detection than with claim verification or world-level reasoning. Across the CD rows, DFDC also remains the hardest benchmark, with even strong methods mostly landing in the mid-70s to low-80s, which further cautions against naive rank ordering.

Method	Layer	Protocol	CDFv2	DFDCP	DFDC
Layer 1					
FreqBlender (Zhou et al., 2024)	L1	CD	94.6	87.6	74.6
SeeABLE (Larue et al., 2023)	L1	CD	87.3	86.3	75.9
LSDA (Yan et al., 2024)	L1	CD	83.0	81.5	73.6
Style Latent Flows (Choi et al., 2024)	L1	CD	89.0	–	–
Layer 2					
TALL-Swin (Xu et al., 2023)	L2	CD	90.8	–	76.8
LipForensics (Haliassos et al., 2021)	L2	CD	82.4	–	73.5
Two-branch (Masi et al., 2020)	L2	CD	76.6	–	–
Layer 3					
MDS (Chugh et al., 2020)	L3	ID	–	–	90.6
Layer 4					
RepDFD (Lin et al., 2025a)	L4	CD	89.9	95.0	81.0
LVLMDFD (Yu et al., 2025)	L4	CD	94.3	92.4	77.0

Table 6 AUC (%) snapshot across the four-layer taxonomy with explicit protocol labels. CD = cross-dataset with training on FF++ and testing on CDFv2, DFDCP, and DFDC; ID = in-domain. CD and ID values are **not directly comparable**, so the table should be read as a protocol-aware snapshot rather than a single ranked summary. “–” denotes not reported.

5 Evaluation of AIGC-V detection

We provide a comprehensive overview of a generic evaluation framework for AIGC-V detection, focusing on evaluation metrics under the dual-view setting and benchmarks organized in line with three AIGC-V paradigms. We also summarize adjacent non-detector-first diagnostic resources for synthetic-video factual-fidelity evaluation.

5.1 Evaluation Metrics

Shared Basic Metrics. We report standard binary metrics for real-versus-AIGC-V detection, including *Acc*, *AUC*, *Precision*, *Recall*, *F1*, *EER*, and *PR-AUC* under class imbalance. Scores may be computed at the frame level or video level, with temporal aggregation such as mean pooling or voting. These metrics provide a shared baseline, but they do not by themselves diagnose temporal coherence, physical plausibility, or semantic consistency.

Visual View Metrics. The visual view evaluates whether a detector captures perceptual evidence that a video *looks and moves* like a real capture. Two aspects matter. **(i) Intrinsic-cue robustness:** acquisition and generation artifacts such as sensor and ISP traces, resampling and coding footprints, and sampling artifacts are still commonly measured with *Acc*, *AUC*, and *EER* (Cheng et al., 2024; Zhou et al., 2024; Gu et al., 2022b), but the decisive protocols are cross-dataset transfer and perturbation stress tests over codec, bitrate, and resolution. Fixed-operating-point metrics such as $TPR@FPR=\alpha$ are also important in deployment-oriented settings. **(ii) Spatiotemporal and physical consistency:** evaluation asks whether motion, layout, and interactions remain temporally and physically plausible (Feng et al., 2023; Zhang et al., 2024a; Anand et al., 2025; Zhang et al., 2025e; Zheng et al., 2025; Kim et al., 2025; Xu et al., 2024). Video-level reporting, temporal-perturbation drops, and motion-ablation studies are generally more informative than strong frame-level scores alone (Kundu et al., 2025; Nie et al., 2024; Yan et al., 2025). Overall, visual-view evaluation should prioritize video-level robustness rather than closed-set frame discrimination.

Language View Metrics. The language view evaluates whether video, audio, and accompanying text or metadata jointly support plausible *facts about the world*. It also decomposes into two aspects. **(i) Cross-modal alignment:** benchmarks probe lip-audio synchrony, speaker identity consistency, caption-video matching, and temporally localized mismatches. Depending on the setting, works combine *Acc* and *AUC* with *AP*, *AR*, *Recall@K*, and *mAP* over *IoU* thresholds, often under modality corruption or degraded-channel tests (Katamneni and Rattani, 2024; Xu et al., 2025c; Wang et al., 2025b; Anshul et al., 2025). **(ii) World knowledge and reasoning:** for fact-level and narrative-level verification, standard classification scores are insufficient. Evaluation therefore adds human judgments, pairwise preferences, question answering, and rationale-quality metrics such as *BLEU*, *ROUGE-L*, *METEOR*, *CIDEr*, and embedding-based similarity (Yu et al., 2025; Shen et al., 2025; Wen et al., 2025b; Sun et al., 2025; Gao et al., 2025; Fu et al., 2025; Hondru et al., 2025). Recent studies further show that observers rely on mixed visual, vocal, and knowledge cues, and that detector outputs can mislead downstream claim verification when they are not tied to explicit evidence (Chen and Goh, 2026; Shi et al., 2026; Sagar et al., 2026). In short,

language-view evaluation shifts the question from “does it look real” to “does it support a plausible, well-grounded story.”

What Evaluation Must Show

- ▶ **AUC is necessary but insufficient:** it measures separability on a protocol, not whether the decision is evidentially well grounded.
- ▶ **Protocols should follow the claimed evidence pathway:** robustness under shift for visual cues, and alignment, localization, grounding, and claim verification with reasoning.
- ▶ **Benchmark numbers are not interchangeable.** In-domain accuracy, cross-dataset transfer, operating-point robustness, and explanation quality answer different questions and should not collapse into one headline score.

5.2 Benchmarks

Benchmark development still follows the three AIGC-V paradigms in Section 2, but the three branches now differ not only in scale, but also in annotation granularity and evidential reach. Taken together, these three paradigm-specific benchmark families define the core detector-evaluation landscape. Table 7 summarizes this landscape of benchmarks for AIGC-V detection. A complementary line of adjacent diagnostic resources, discussed later in this subsection, broadens that landscape beyond detector-first evaluation toward physical-rule testing, world-dynamics and causality probes, and explanation-oriented diagnosis.

LMV Related Benchmarks. LMV related benchmarks remain the historical backbone of AIGC-V detection because partial edits preserve most of the source video while exposing localized forensic traces. That structure makes LMV especially suitable for testing artifact sensitivity, compression robustness, codec variation, and cross-dataset transfer. Among the foundational resources, FaceForensics++ (Rossler et al., 2019) was decisive because it standardized post-compression manipulation protocols and effectively set the early evaluation culture; Celeb-DF (Li et al., 2020b) and DFDC (Dolhansky et al., 2020) then pushed the field toward harder realism and larger scale; DeeperForensics-1.0 (Jiang et al., 2020) made robustness under real-world distortions a central concern rather than an afterthought. ForgeryNet (He et al., 2021) broadened the manipulation space further. Later resources pushed LMV in more deployment-facing directions: WildDeepfake (Zi et al., 2020) and KoDF (Kwon et al., 2021) moved toward in-the-wild and cross-lingual conditions, AI-Face (Lin et al., 2025b) enlarged demographic coverage, and CDDDB (Li et al., 2023) together with DeepfakeBench (Yan et al., 2023) reframed LMV as a protocol problem rather than a single-dataset problem. DD-VQA (Zhang et al., 2024b) and ExDDV (Hondru et al., 2025) begin to attach question answering and explanation supervision to a regime that was originally almost purely artifact-driven, while Beyond Static Artifacts (Gu et al., 2026) turns temporal deepfake reasoning itself into an explicit benchmark target for VLMs.

AVE Related Benchmarks. AVE related benchmarks are organized around whether speech, mouth motion, speaker identity, and language-bearing cues remain jointly coherent over time, so they are less about static artifact capture than about temporal and cross-modal verification. This branch remains smaller than LMV partly because realistic audio-visual edits require synchronized manipulation across modalities together with finer temporal

annotation. FakeAVCeleb (Khalid et al., 2021) established the basic paired audio-video manipulation setting, but LAV-DF (Cai et al., 2022) was especially important because it turned localized, content-driven audio-visual mismatch into a first-class benchmark target rather than treating the clip as uniformly fake. Scale then increased with FakeMix (Jung et al., 2024) and AV-Deepfake1M (Cai et al., 2024). The more recent resources differentiate the evaluation question itself: DigiFakeAV (Liu et al., 2025b) targets diffusion-based digital humans, VCapAV (Wang et al., 2025e) matters because it converts AVE into a caption-grounded inconsistency problem, MAVOS-DD (Croitoru et al., 2025) studies multilingual open-set conditions, X-AVFake (Chen et al., 2025a) introduces grounded language reasoning, and MMDF (Kim et al., 2026) broadens manipulation and generator coverage. Evaluation in this branch therefore centers on lip-audio alignment, speaker-content consistency, dubbed or spliced segments, temporal localization, and robustness under degraded settings.

GVS Related Benchmarks. GVS related benchmarks are the fastest-moving branch because fully generated videos weaken explicit editing traces while amplifying generator diversity, plausibility failures, and transfer risk. Early resources, including GVF (Ma et al., 2024a), GenVidDet (Ji et al., 2024), GenVideo (Chen et al., 2024a), DVF (Song et al., 2024b), and GenVidBench (Ni et al., 2025), established fully generated video detection as a distinct benchmark regime rather than an extension of face-forgery evaluation. Later benchmarks broadened both scope and supervision: LOKI (Ye et al., 2025) and Deepfake-Eval-2024 (Chandra et al., 2025) expanded modality coverage and in-the-wild conditions; GenWorld (Chen et al., 2025c) shifted attention toward real-world simulation; DAVID-X (Gao et al., 2025) and DeeptraceReward (Fu et al., 2025) introduced defect-level or human-perceived trace supervision; and ER-FF++set (Sun et al., 2025) together with AEGIS (Li et al., 2025a) continued the move toward explanation-aware evaluation. The newest large-scale suites, including AIGVDBench (Ma et al., 2026), SynthForensics (Leotta et al., 2026), ViFBench (Li et al., 2025b), Video Reality Test (Wang et al., 2025c), and MintVid (Tan et al., 2026), emphasize generator diversity, stronger realism, and deployment-oriented stress testing. AIGVDBench is particularly consequential because its scale makes cross-generator transfer patterns interpretable at ecosystem level. As a result, GVS benchmarks increasingly evaluate cross-generator transfer, physical plausibility, event logic, and semantically fabricated scenarios rather than artifact sensitivity alone.

What Benchmarks Actually Cover

- ▶ **Coverage is structurally uneven.** LMV has the deepest protocol history, AVE remains narrower because synchronized multimodal annotation is costly, and GVS changes fastest as generator turnover keeps resetting the task.
- ▶ **The three paradigms test different evidence pathways.** LMV emphasizes localized forensic residue and transfer, AVE temporal and cross-modal coherence, and GVS transfer under weaker artifacts together with plausibility and event logic.
- ▶ **Benchmark difficulty shifts with the generation paradigm.** LMV mainly asks whether traces survive compression and domain shift, AVE whether inconsistencies can be localized across modalities, and GVS whether detectors remain reliable as generators become more realistic and diverse.

Benchmark & Dataset	Description	Paradigms	Date
A. Local Manipulation Video (LMV)			
FaceForensics++ (Rossler et al., 2019)	Forensics dataset with 1,000 original videos.	LMV	01/2019
Celeb-DF (Li et al., 2020b)	Large-scale challenging deepfake-forensics dataset.	LMV	09/2019
DeeperForensics-1.0 (Jiang et al., 2020)	Large-scale dataset for real-world face-forgery detection.	LMV	05/2020
DFDC (Dolhansky et al., 2020)	Large-scale face-swapping dataset (mostly local forgeries).	LMV	06/2020
WildDeepfake (Zi et al., 2020)	In-the-wild dataset for AIGC-V detection.	LMV	10/2020
ForgeryNet (He et al., 2021)	Mega-scale benchmark for image and video face-forgery analysis.	LMV	06/2021
KoDF (Kwon et al., 2021)	Large-scale Korean AIGC-V detection dataset.	LMV	10/2021
CDDb (Li et al., 2023)	Benchmark for easy, hard, and long-sequence AIGC-V detection.	LMV	01/2023
DF-Platter (Narayan et al., 2023)	Multi-face heterogeneous deepfake dataset.	LMV	06/2023
DeepfakeBench (Yan et al., 2023)	Comprehensive AIGC-V detection benchmark.	LMV	07/2023
AI-Face (Lin et al., 2025b)	Million-scale, demographically annotated AI-face dataset.	LMV&GVS	06/2024
DD-VQA (Zhang et al., 2024b)	AIGC-V detection VQA (image-question-answer triplets).	LMV	09/2024
ExDDV (Hondru et al., 2025)	Benchmark for explainable AIGC-V video detection.	LMV	11/2025
FAQ / Beyond Static Artifacts (Gu et al., 2026)	Forensic benchmark and instruction-tuning resource for temporal video deepfake reasoning in VLMs.	LMV	02/2026
B. Audio-Visual Editing (AVE)			
FakeAVCeleb (Khalid et al., 2021)	Audio-visual multimodal deepfake dataset.	AVE	08/2021
LAV-DF (Cai et al., 2022)	Content-driven localized audio-visual deepfake dataset.	AVE	11/2022
FakeMix (Jung et al., 2024)	Clip-level benchmark for manipulated audio and video segments.	AVE	08/2024
AV-Deepfake1M (Cai et al., 2024)	Large-scale LLM-driven audio-visual deepfake dataset.	AVE	10/2024
ArEnAV (Kuckreja et al., 2025)	Audio-visual deepfake dataset for Arabic-English code-switching (CSW).	AVE	05/2025
MAVOS-DD (Croitoru et al., 2025)	Multilingual open-set benchmark for audio-visual AIGC-V detection.	AVE	05/2025
DigiFakeAV (Liu et al., 2025b)	Benchmark for diffusion-based digital-human audio-visual forgeries.	AVE	05/2025
SocialDF (Batra et al., 2025)	2,126 social-media videos with real and deepfake content under SOTA manipulations.	AVE	07/2025
VCapAV (Wang et al., 2025e)	Video-caption-based audio-visual AIGC-V detection dataset.	AVE	08/2025
X-AVFake (Chen et al., 2025a)	Dual-modality manipulations with grounded language reasoning.	AVE	08/2025
AV-Deepfake1M++ (Cai et al., 2025)	2M clips extending AV-Deepfake1M with diverse audio-visual manipulations.	AVE	10/2025
MMDF (Kim et al., 2026)	Multimodal deepfake dataset spanning GAN, diffusion, and flow-matching manipulations.	AVE	03/2026
C. Generative Video Synthesis (GVS)			
GVF (Ma et al., 2024a)	Generated Video Forensics benchmark for AI-video detection.	GVS	02/2024
GenVidDet (Ji et al., 2024)	Real and AIGC-V videos from eight generation models.	GVS	05/2024
GenVideo (Chen et al., 2024a)	AIGC-V detection dataset.	GVS	05/2024
DVF (Song et al., 2024b)	Diffusion Video Forensics dataset and benchmark.	GVS	12/2024
GenVidBench (Ni et al., 2025)	AIGC-V detection dataset.	GVS	01/2025
Deepfake-Eval-2024 (Chandra et al., 2025)	Multi-modal in-the-wild benchmark of deepfakes from 2024.	L&A&G	03/2025
LOKI (Ye et al., 2025)	Multimodal synthetic-data detection benchmark over video, image, 3D, text, and audio.	GVS	04/2025
GenBuster-200K (Wen et al., 2025a)	Real + synthetic videos simulating real-world conditions.	GVS	05/2025
GenWorld (Chen et al., 2025c)	Large-scale real-world simulation dataset for AI-video detection.	GVS	06/2025
Ivy-Fake (Jiang et al., 2025)	Large-scale benchmark for explainable AIGC detection.	GVS	06/2025
DAVID-X (Gao et al., 2025)	AIGC-V with defect-level spatiotemporal annotations + rationales.	GVS	06/2025
GenBuster++ (Wen et al., 2025b)	A cross-modal benchmark for VLM evaluation.	GVS	07/2025
DeeptraceReward (Fu et al., 2025)	Spatiotemporal benchmark with human-perceived fake traces.	GVS	09/2025
ER-FF++set (Sun et al., 2025)	Benchmark with supervision for detection + explanation.	GVS	10/2025
AEGIS (Li et al., 2025a)	Large-scale benchmark for AIGC-V authenticity.	GVS	10/2025
ViFBench (Li et al., 2025b)	3K videos from 10+ SOTA generators.	GVS	12/2025
Video Reality Test (Wang et al., 2025c)	An ASMR-sourced benchmark.	GVS	12/2025
AIGVDBench (Ma et al., 2026)	Large-scale benchmark with 440K videos from 31 generation models and 33 evaluated detectors.	GVS	01/2026
SynthForensics (Leotta et al., 2026)	Human-centric synthetic-video benchmark with 6,815 videos from five open-source T2V generators.	GVS	02/2026
MintVid (Tan et al., 2026)	Lightweight high-quality benchmark with 3K videos from nine state-of-the-art generators.	GVS	02/2026

Table 7 Overview of existing datasets and benchmarks through March 2026, and their alignment with the AIGC-V paradigms introduced in our survey.

Diagnostic Resources Beyond Detection Benchmarks. Adjacent diagnostic resources extend evaluation along a clear progression: from basic rule violations, to coherent world dynamics, and finally to explanation-oriented diagnosis. Physical Rule Violations asks whether generated videos obey basic physical constraints: VideoPhy (Bansal et al., 2025), Physics-IQ (Motamed et al., 2025b), and IPV-Bench (Bai et al., 2025) stress physical commonsense and impossible scenarios; Morpheus (Zhang et al., 2025a) and T2VPhysBench (Guo et al., 2025) tighten this into experiment-grounded and first-principles testing; and PhyWorldBench (Gu et al., 2025), VideoPhy-2 (Bansal et al., 2026), and Physion-Eval (Zhang et al., 2026) move further toward action-centric settings, localized glitches, and expert reasoning traces. World Dynamics and Causality then asks whether those local regularities scale to temporally coherent events and world knowledge over time. WorldSimBench (Qin et al., 2025) makes the world-simulator question explicit; PhyGenBench (Meng et al., 2025) studies multi-step intuitive physics; StoryEval (Wang et al., 2024c) turns consecutive events into a story-level stress test; VideoVerse (Wang et al., 2025f) and T2VWorldBench (Chen et al., 2025d) probe temporal causality, hidden semantics, and world knowledge; and SVBench (Peng et al., 2025) extends the agenda to socially coherent behavior. Controlled analyses (Kang et al., 2024) further argue that this world-model framing should remain a benchmark target rather than an assumed capability. Explanation-Oriented Diagnosis then asks whether a system can turn those failures into explicit, interpretable judgments. SPOTLIGHT (Chinchure et al., 2025) and VideoHallu (Li et al., 2025c) emphasize error localization and hallucination-aware synthetic-video understanding, while TRAVL (Motamed et al., 2025a) and PhyDetEx (Wang et al., 2025g) treat implausibility judgment and explanation as trainable capabilities. Table 8 summarizes this adjacent landscape.

Work and Resource	Description	Type	Date
A. Physical Rule Violations			
VideoPhy (Bansal et al., 2025)	Benchmarks whether generated videos satisfy physical commonsense about objects, attributes, and interactions.	Eval	06/2024
Physics-IQ (Motamed et al., 2025b)	Probes whether text-to-video models obey basic physical principles under controlled prompts.	Eval	01/2025
IPV-Bench (Impossible Videos) (Bai et al., 2025)	Designs impossible scenarios across physical, geographical, biological, and social domains for stress testing generated videos.	Eval	03/2025
Morpheus (Zhang et al., 2025a)	Benchmarks physical reasoning of video generative models with real physical experiments and measurable conservation-law violations.	Eval	04/2025
T2VPhysBench (Guo et al., 2025)	Tests first-principles physical consistency in text-to-video generation, including counterfactual robustness and state-transition fidelity.	Eval	05/2025
PhyWorldBench (Gu et al., 2025)	Evaluates physical realism of text-to-video models with physics-grounded prompts and anti-physics stress cases.	Eval	07/2025
VideoPhy-2 (Bansal et al., 2026)	Extends physical-commonsense evaluation to more challenging action-centric and interaction-heavy generated videos.	Eval	01/2026
Physion-Eval (Zhang et al., 2026)	Provides expert reasoning traces, localized physical glitches, and explanations for physical-realism failures in generated videos.	Eval	03/2026
B. World Dynamics and Causality			
WorldSimBench (Qin et al., 2025)	Evaluates whether video generators behave like world simulators by combining perceptual quality and embodied action consistency.	Eval	10/2024
Towards World Simulator (PhyGenBench) (Meng et al., 2025)	Crafts a physical-commonsense benchmark for video generation with multi-step dynamics and simulator-style diagnostics.	Eval	10/2024
StoryEval (Wang et al., 2024c)	Benchmarks whether T2V models can present short stories composed of consecutive events for future long-video generation.	Eval	12/2024
T2VWorldBench (Chen et al., 2025d)	Evaluates world-knowledge generation across physics, nature, activity, culture, causality, and object domains.	Eval	07/2025
VideoVerse (Wang et al., 2025f)	World-model-oriented benchmark with hidden semantics, event-level temporal causality, and world-knowledge questions.	Eval	10/2025
SVBench (Peng et al., 2025)	Evaluates social reasoning in video generation, including intention, joint attention, norms, and prosocial behavior.	Eval	12/2025
C. Explanation-Oriented Diagnosis			
TRAVL (Motamed et al., 2025a)	Adds intra-frame spatial and trajectory-guided temporal attention and fine-tunes on ImplausiBench to improve VLM judgments of physically implausible videos.	Method	10/2025
SPOTLIGHT (Chinchure et al., 2025)	Fine-grained identification and localization of generation errors using VLMs.	Eval	11/2025
VideoHallu (Li et al., 2025c)	Evaluates multimodal hallucination on synthetic videos and studies mitigation via targeted post-training.	Eval	12/2025
PhyDetEx (Wang et al., 2025g)	Introduces a benchmark and trains models to detect and explain violated physical rules in text-to-video outputs.	Method	12/2025

Table 8 Adjacent factual-fidelity diagnostic resources relevant to AIGC-V detector design and world-model-oriented evaluation.

6 Challenges and Future

6.1 Robust Diagnostic Evaluation

Chen et al. (2024b) and Sun et al. (2025) argue that classification-centric metrics such as clip-level *AUC/EER* are often not sufficiently evidential or interpretable for security-sensitive settings. Under our *factual-fidelity* objective in §3.1, strong closed-set discrimination does not necessarily mean the detector can substantiate *which* real-world proposition is violated, *where* the violation occurs, or *why* the judgment is reliable. Evaluation should therefore be *diagnostic*: it should probe the detector’s reliance on complementary pathways (perceptual/temporal cues vs. claim-level semantic verification), and reveal the failure modes that lead to factual-fidelity breakdown.

Concretely, one critical axis is robustness of perceptual and temporal evidence under distribution shift. As diffusion and large video models increasingly suppress noticeable per-frame artifacts, evaluation must stress-test video-level consistency and robustness-to-shift rather than only in-domain performance. This pressure is already visible in explainable detection settings such as EDVD-LLaMA (Sun et al., 2025) and DeeptraceReward (Fu et al., 2025), in generator-diverse detection benchmarks such as GenVidBench (Ni et al., 2025), and even in physics-oriented generative evaluation such as Physics-IQ (Motamed et al., 2025b). Serrano et al. (2026) and Hasan et al. (2026) further show that strong clean-set scores can mask severe brittleness under transfer-based adversarial attacks and modern synthesis pipelines, leaving a wide gap between benchmark discrimination and deployment robustness. This motivates cross-dataset generalization protocols, systematic perturbation sweeps over codec, bitrate, and resolution, and deployment-oriented reporting at fixed operating points such as $TPR@FPR=\alpha$ to better reflect real risk.

A second axis is whether the system can treat a video as making checkable claims and verify them with grounded evidence. Beyond synchronization, cross-modal alignment tests for retrieval and temporal localization should be complemented with fact-level and knowledge-level verification. Common sense reasoning for deepfake detection (Zhang et al., 2024b) pushes evaluation toward explicit question answering, LVLM-DFD (Yu et al., 2025) and AuthGuard (Shen et al., 2025) emphasize grounded multimodal reasoning, and BusterX++ (Wen et al., 2025b) adds unified cross-modal detection-and-explanation settings. These developments also motivate measuring rationale quality and grounding, not just label accuracy. Together, these diagnostics align evaluation with the shift from perceptual-fidelity discrimination to factual-fidelity verification.

6.2 Claim-Level and Dynamic Evaluation

A natural next step is to make benchmarks *evidence-first* by moving from clip-level labels to claim-level supervision. In this setting, each clip is decomposed into a compact set of propositions specifying who, when, where, and what happens, and each proposition is paired with timestamped in-video evidence. Evaluation then scores both claim correctness and evidence grounding, making “what is fake” and “where it is fake” traceable rather than implicit. This direction is consistent with recent explainable AIGC detection resources

that augment binary supervision with rationales, traces, and defect-level annotations. X2-DFD (Chen et al., 2024b) and EDVD-LLaMA (Sun et al., 2025) make explanation part of the detector output, Ivy-Fake (Jiang et al., 2025) and DAVID-XR1 (Gao et al., 2025) attach rationales and defect-level traces to AI-generated video analysis, DeepttraceReward (Fu et al., 2025) adds human-perceived fake traces, and ExDDV (Hondru et al., 2025) turns explainable supervision into a benchmark design target.

To prevent shortcut learning and isolate genuine reasoning, evidence-first benchmarks should also incorporate targeted stress tests. One complementary approach is claim-level stress testing for event logic: construct counterfactual or adversarial cases by rewriting event scripts or state transitions while controlling visual-quality confounds, so performance differences primarily reflect relational and event reasoning rather than superficial appearance cues. This concern is explicit in Can We Leave Deepfake Data Behind in Training Deepfake Detector? (Cheng et al., 2024) and FakeChain (Heo and Woo, 2025), both of which expose how easily detectors can rely on shallow cues. In addition, benchmarks increasingly include supervision for tampered locations, spatiotemporal traces, and grounded explanations. EDVD-LLaMA (Sun et al., 2025), DeepttraceReward (Fu et al., 2025), ExDDV (Hondru et al., 2025), and DAVID-XR1 (Gao et al., 2025) all move beyond binary labels in this direction, while Video Reality Test (Wang et al., 2025c) adds realistic high-fidelity content where obvious artifacts are weak. Where available, cross-validating content-side decisions with provenance and authentication signals can further improve trustworthiness in security-sensitive settings.

Finally, given fast-evolving generators, we expect evaluation to move beyond static test sets toward a closed loop between a *real-time detection platform* and a *dynamic benchmark*. In an arena- or leaderboard-style regime, the benchmark is continuously refreshed with outputs from newly released generators and with realistic ingestion or transcoding effects such as codec and resolution changes; detectors are then periodically re-evaluated under a unified protocol to produce regression trajectories rather than one-off scores (Wang et al., 2025c). Such a setup measures not only snapshot accuracy, but also robustness decay, adaptation lag, and failure modes under sustained distribution shift. Hard cases that repeatedly fool strong detectors can then be promoted into subsequent benchmark refreshes with failure-type annotations, turning deployment feedback into iterative stress testing rather than passive reporting.

6.3 Unified Explainable Detection

Wang et al. (2025c) indicates that, for AIGC-V with sparse factual content but high perceptual fidelity and well-aligned modalities, VLMs can perform near random in perceptual-fidelity discrimination, far below humans. This highlights a core gap of language-only detection and the need for visual-view evidence. At the same time, stronger generators and anti-detection pipelines make it increasingly unsafe to assume that authenticity can be decided from perceptual traces alone. Common sense reasoning (Zhang et al., 2024b), generalizable LVLM reasoning (Yu et al., 2025), language guidance (Shen et al., 2025), explainable video reasoning (Gao et al., 2025), and memory-anchored multimodal reasoning (Chen et al., 2025a) all point toward detectors that must reason over semantic and factual content

as well. When perceptual fidelity is high and modalities align, decisive evidence may still come from subtle visual statistics or long-range temporal inconsistencies; frequency-aware blending cues (Zhou et al., 2024), local dynamic inconsistency learning (Gu et al., 2022a), natural consistency representations (Zhang et al., 2024a), and plug-and-play video-level blending with spatiotemporal adapter tuning (Yan et al., 2025) all support this view. Conversely, when generators suppress artifacts and preserve visual coherence, visual-only cues can be insufficient and fact-level reasoning becomes necessary; this is exactly the trajectory emphasized by common-sense and reasoning-heavy detection methods (Zhang et al., 2024b; Yu et al., 2025; Shen et al., 2025; Gao et al., 2025). Unified detection is therefore a two-pathway problem: perceptual evidence plus fact-level verification.

In practice, this motivates cross-layer evidence fusion: frame-level intrinsic cues from Layer 1, sequence-level motion and physics coherence from Layer 2, multimodal verification from Layer 3, and external-knowledge reasoning from Layer 4 should be integrated into a unified evidence graph, where corroborating evidence boosts confidence and conflicts trigger abstention or human review. The goal, however, is not to average heterogeneous scores, but to preserve evidential granularity so that low-level traces, localized cross-modal conflicts, and claim-level checks remain individually inspectable. The key challenge is to connect low-level cues to high-level conclusions in an auditable way, as X2-DFD (Chen et al., 2024b) and EDVD-LLaMA (Sun et al., 2025) already make clear. A practical system should output structured evidence objects, such as suspicious segments and defect-level traces (Gao et al., 2025; Anshul et al., 2025), localized cross-modal mismatches (Katamneni and Rattani, 2024; Chen et al., 2024b), and tested claims or entities (Zhang et al., 2024b), rather than only free-form rationales. Tool-augmented pipelines such as LAVID (Liu et al., 2025d) help only when each tool invocation is tied to a concrete sub-claim and yields reusable evidence objects; memory-anchored multimodal reasoning (Chen et al., 2025a) points to the same requirement.

6.4 Evidence-First Trustworthy Detection

To match trustworthy detection that truly targets factual fidelity, system design should follow an Evidence-First principle with an explicit reasoning path: Identify, Localize, and Explain. X2-DFD (Chen et al., 2024b), DeeptraceReward (Fu et al., 2025), and EDVD-LLaMA (Sun et al., 2025) already move in this direction. The key implication is that explanation should be downstream of evidence extraction, not a post-hoc verbal gloss over a classifier score. This path should separate candidate-issue discovery, where-and-when localization, and explanation of which perceptual, temporal, cross-modal, or fact-level constraints are violated. Such a design requires consistent evidence schemas, as DAVID-XR1 (Gao et al., 2025) and X2-DFD (Chen et al., 2024b) suggest, together with calibrated uncertainty emphasized by reliability-centered analyses (Wang et al., 2024a). Calibration matters because trustworthy deployment depends on knowing when evidence is weak, conflicting, or incomplete rather than forcing a binary answer.

Within this dual-pathway setup, each tool invocation or knowledge citation should be bound to a specific argumentation step (Liu et al., 2025d). Content-side analysis should also be cross-validated against source-side provenance and authentication signals when

available, rather than treated as an unrelated defense channel. This connection is explicit in media-forensics and defense surveys (Verdoliva, 2020; Deng et al., 2025), in technical provenance standards such as C2PA (Coalition for Content Provenance and Authenticity, 2024), and in complementary verification frameworks built on face geometry (Tursman et al., 2020) or blockchain-backed authenticity pipelines (Hajjej et al., 2026). This matters beyond standalone detection accuracy: Sagar et al. (2026) suggests that detector outputs can become misleading priors when they are injected into claim verification pipelines without explicit evidence grounding. The practical difficulty is to reconcile potentially conflicting evidence and surface it in a unified evidence space that supports auditing and abstention under uncertainty (Wang et al., 2024a). A trustworthy system should therefore preserve both agreement and disagreement across evidence sources, rather than collapsing them too early into a single confidence score.

What Trustworthy Detection Requires

- ▶ **Diagnose the real failure.** Evaluation should reveal where factual fidelity breaks, not only whether real and fake clips separate on a closed set.
- ▶ **Supervise evidence, not only labels.** Claim-level annotations, timestamped evidence, short-cut controls, and dynamic testbeds are needed as generators evolve.
- ▶ **Return structured evidence, not only scores.** Cross-layer fusion is useful only when cues become traceable claims, segments, entities, or mismatches rather than free-form confidences or rationales.
- ▶ **Treat deployment as evidence governance.** Provenance checks, calibration, and abstention are part of trustworthy detection when evidence conflicts or remains incomplete.

7 Related Work

Table 9 systematically compares representative surveys by what they take as the task, what deployment concerns they foreground, and how far they move from artifact-centric detection toward grounded verification. This comparison clarifies that our contribution is not only broader coverage, but also a different organizing principle for AIGC-V detection.

7.1 Scope and Framing

Video-targeted scope. Most prior surveys take either a face-manipulation view or a broad *deepfake* or AIGC scope, so they naturally organize the field around generation pipelines, detector families, and generic forensic cues. This breadth is useful for coverage, but it often compresses important distinctions among local manipulation, audio-visual editing, and fully generated video, even though these regimes differ in threat model, evidence availability, and evaluation difficulty. Broad overviews by Tolosana et al. (2020), Mirsky and Lee (2021), and Nguyen et al. (2022) are representative of this framing. From a media-forensics viewpoint, Verdoliva (2020) places deepfakes within a larger integrity-verification pipeline of detection, localization, attribution, and authentication. More video-oriented surveys, including Hashmi et al. (2024) on audio-visual deepfakes and Kaur et al. (2024) on video-specific challenges, move closer to our focus, but they still do not systematize the full AIGC-V landscape across the three paradigms used here.

Survey	Scope & Framing			Language View		
	Video-Targeted Scope	Trust-worthiness Framing	Factual-Fidelity Verif.	Explain-ability & Loc.	Semantic Cues	World-level Reasoning
Wang et al. (2024a)	✗	✓	✗	✗	✗	✗
Xu et al. (2025a)	✗	✓	✗	✗	✗	✗
Croitoru et al. (2024)	✗	✓	✗	✗	✗	✗
Hashmi et al. (2024)	✓	✗	✗	✗	✗	✗
Kaur et al. (2024)	✓	✓	✗	✗	✗	✗
Deng et al. (2025)	✗	✓	✗	✓	✗	✗
Nguyen-Le et al. (2025)	✗	✓	✗	✓	✗	✗
Lin et al. (2024a)	✗	✓	✗	✓	✓	✗
Zou et al. (2025)	✗	✓	✗	✓	✓	✗
Ours	✓	✓	✓	✓	✓	✓

Table 9 Comparison of representative surveys and ours across (i) **Scope and Framing** (video-targeted scope; trustworthiness framing; factual-fidelity verification) and (ii) **Language View** (explainability and localization; semantic cues; world-level reasoning). ✓ indicates substantial discussion; ✗ indicates absent or brief mention.

Trustworthiness framing. Recent surveys increasingly move beyond raw accuracy. Wang et al. (2024a) and Deng et al. (2025) shift attention toward deployment reliability, adversarial evolution, uncertainty, failure modes, and complementary defenses such as authentication and disruption. Beyond the deepfake-only scope, Lin et al. (2024a) surveys AI-generated multimedia across modalities and explicitly separates better classification from broader goals such as generalizability, robustness, and interpretability. Zou et al. (2025) further highlights the transition from domain-specific detectors to VLM- and MLLM-based systems with localization and explanation. Yet even under this framing, trustworthiness is still usually defined around making detector outputs more reliable, rather than around making the underlying authenticity judgment explicitly evidential and claim grounded.

Factual-fidelity verification. Accordingly, “verification” in prior surveys usually appears either inside a broader media-forensics workflow (Verdoliva, 2020) or inside a deployment-reliability discussion centered on robustness, uncertainty, and countermeasures (Wang et al., 2024a; Deng et al., 2025). Both perspectives are important, but neither reorganizes the task around what a video claims and whether those claims remain consistent with the world. Lin et al. (2024a) and Zou et al. (2025) bring in semantic cues and interpretability with localization, but the emphasis still falls on recognizing manipulations, localizing suspicious regions, or justifying predictions. Fact-level checking of who did what, where, and when is rarely treated as the primary unit of analysis. As summarized in Table 9, this is the main framing gap that motivates our survey.

7.2 Language View

Explainability & localization. Recent surveys increasingly recognize explainability and localization, especially in trustworthy and VLM-based detection. Lin et al. (2024a), Zou

et al. (2025), Deng et al. (2025), and Nguyen-Le et al. (2025) all acknowledge that deployment requires more than a scalar real/fake score. However, explanation is still usually understood as saliency, manipulated-region localization, or natural-language justification attached to a prediction. What remains less explicit is whether those outputs become reusable evidence objects that can support later claim verification and auditing.

Semantic cues. As deepfakes move from static artifacts to temporally coherent synthesis, surveys increasingly emphasize video-level evidence and multimodal cues beyond pixels. Kaur et al. (2024) highlights video-specific challenges and robustness under real-world degradations, while Hashmi et al. (2024) centers audio-visual correspondence, identity preservation, and synchrony cues. Lin et al. (2024a) and Zou et al. (2025) push further by acknowledging semantic cues in language-enabled detectors, and Nguyen-Le et al. (2025) together with Xu et al. (2025a) consolidate passive detection cues across modalities and emerging AIGC forensics directions. Even so, semantic information is usually introduced as auxiliary context that improves classification or explanation, not as a primary evidential pathway for testing factual fidelity.

World-level reasoning. Explicit world-level reasoning about events, causality, and plausibility remains even less developed in the survey literature. Most prior surveys stop at semantic alignment or multimodal consistency and do not elevate world knowledge, event logic, or external verification into a first-class evidence pathway. Table 9 makes this asymmetry clear: semantic cues are beginning to appear, but world-level reasoning and, especially, fact-level factual-fidelity verification remain under-emphasized. Our survey addresses this gap by adopting a video-targeted scope, foregrounding factual-fidelity verification, and organizing methods, evaluation, and benchmarks under our taxonomy.

8 Conclusion

This survey reframes AI-generated video detection as factual-fidelity verification and synthesizes prior work through a Vision-Language Dual-View four-layer taxonomy with aligned benchmarks and metrics. Beyond summarizing methods, we connect these layers to deployment challenges, evaluation gaps, and emerging trends, and distill key requirements for trustworthy detection, including evidence-first and traceable decisions as well as robustness across generators and real-world scenarios. We expect this survey to provide a clear and actionable reference for future AIGC-V detection research and practice.

More broadly, trustworthy AIGC-V detection should become a shared agenda for the computer vision, natural language processing, multimodal, and world model communities. CV contributes spatiotemporal evidence and forensic robustness; NLP contributes decomposition, reasoning, grounding, and explanation; multimodal and world model research contribute stronger cross-modal alignment and richer priors about physical, causal, and temporal coherence. Together, these ingredients can push detection beyond artifact hunting toward a stricter notion of realism: not whether a video merely looks plausible, but whether its entities, events, and dynamics remain faithful to the constraints of the real world.

References

- Shruti Agarwal and Hany Farid. Detecting deep-fake videos from aural and oral dynamics. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, pages 981–989, June 2021.
- Shruti Agarwal, Hany Farid, Tarek El-Gaaly, and Ser-Nam Lim. Detecting deep-fake videos from appearance and behavior. In *2020 IEEE international workshop on information forensics and security (WIFS)*, pages 1–6. IEEE, 2020.
- Tharun Anand, Siva Sankar, and Pravin Nair. Detecting localized deepfake manipulations using action unit-guided video representations. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 4341–4351, 2025.
- Ashutosh Anshul, Shreyas Gopal, Deepu Rajan, and Eng Siong Chng. Next-frame feature prediction for multimodal deepfake detection and temporal localization. *arXiv preprint arXiv:2511.10212*, 2025.
- Ashutosh Anshul, Eng Chng, and Deepu Rajan. A-v representation learning via audio shift prediction for multimodal deepfake detection and temporal localization. In *2026 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2026. <https://wacv.thecvf.com/Conferences/2026/AcceptedPapers>. Poster session.
- Marcella Astrid, Enjie Ghorbel, and Djamila Aouada. Audio-visual deepfake detection with local temporal inconsistencies, 2025. <https://arxiv.org/abs/2501.08137>.
- Sujata Bahadure and Vanita Mane. Deepfake detection using deep convolutional neural network and long short-term memory. *Multimedia Tools and Applications*, 85(259), 2026. doi: 10.1007/s11042-026-21375-7. <https://doi.org/10.1007/s11042-026-21375-7>.
- Zechen Bai, Hai Ci, and Mike Zheng Shou. Impossible videos. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 2458–2483. PMLR, 13–19 Jul 2025. <https://proceedings.mlr.press/v267/bai25a.html>.
- Hritik Bansal, Zongyu Lin, Tianyi Xie, Zeshun Zong, Michal Yarom, Yonatan Bitton, Chenfanfu Jiang, Yizhou Sun, Kai-Wei Chang, and Aditya Grover. Videophy: Evaluating physical commonsense for video generation. In *The Thirteenth International Conference on Learning Representations*, 2025. <https://openreview.net/forum?id=5MWhU3t7WR>.
- Hritik Bansal, Clark Peng, Yonatan Bitton, Roman Goldenberg, Aditya Grover, and Kai-Wei Chang. Videophy-2: A challenging action-centric physical commonsense evaluation in video generation. In *The Fourteenth International Conference on Learning Representations*, 2026. <https://openreview.net/forum?id=P8vQYmq3TB>.
- Arnesh Batra, Jashn Khemani, Arush Gumber, Anushk Kumar, Arhan Jain, and Somil Gupta. Socialdf: Benchmark dataset and detection model for mitigating harmful deep-fake content on social media platforms. In *Proceedings of the 4th ACM International Workshop on Multimedia AI against Disinformation*, pages 81–89, 2025.

- Matyas Bohacek and Hany Farid. Lost in translation: Lip-sync deepfake detection from audio-video mismatch. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, pages 4315–4323, June 2024.
- Gaetan Brison, Soobash Daiboo, Samy Aimeur, Awais Hussain Sani, Xi Wang, Gianni Franchi, and Vicky Kalogeiton. Fakeparts: a new family of ai-generated deepfakes. *arXiv preprint arXiv:2508.21052*, 2025.
- ByteDance Seed. Official launch of seedance 2.0. https://seed.bytedance.com/en/seedance2_0, 2026. Accessed: 2026-02-24.
- Zhixi Cai, Kalin Stefanov, Abhinav Dhall, and Munawar Hayat. Do you really mean that? content driven audio-visual deepfake dataset and multimodal method for temporal forgery localization. In *2022 International Conference on Digital Image Computing: Techniques and Applications (DICTA)*, pages 1–10. IEEE, 2022.
- Zhixi Cai, Shreya Ghosh, Aman Pankaj Adatia, Munawar Hayat, Abhinav Dhall, Tom Gedeon, and Kalin Stefanov. Av-deepfake1m: A large-scale llm-driven audio-visual deepfake dataset. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 7414–7423, 2024.
- Zhixi Cai, Kartik Kuckreja, Shreya Ghosh, Akanksha Chuchra, Muhammad Haris Khan, Usman Tariq, Tom Gedeon, and Abhinav Dhall. Av-deepfake1m++: A large-scale audio-visual deepfake benchmark with real-world perturbations. In *Proceedings of the 33rd ACM International Conference on Multimedia*, pages 13686–13691, 2025.
- Songjun Cao, Yuqi Li, Yunpeng Luo, Jianjun Yin, and Long Ma. Leveraging large multimodal models for audio-video deepfake detection: A pilot study. In *2026 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2026. doi: 10.48550/ARXIV.2602.23393. <https://doi.org/10.48550/ARXIV.2602.23393>. Accepted as conference paper.
- Nuria Alina Chandra, Ryan Murtfeldt, Lin Qiu, Arnab Karmakar, Hannah Lee, Emmanuel Tanumihardja, Kevin Farhat, Ben Caffee, Sejin Paik, Changyeon Lee, et al. Deepfake-eval-2024: A multi-modal in-the-wild benchmark of deepfakes circulated in 2024. *arXiv preprint arXiv:2503.02857*, 2025.
- Chen Chen and Dion Hoe-Lian Goh. Seeing, hearing, and knowing together: Multimodal strategies in deepfake videos detection. *CoRR*, abs/2602.01284, 2026. doi: 10.48550/ARXIV.2602.01284. <https://doi.org/10.48550/ARXIV.2602.01284>.
- Chen Chen, Runze Li, Zejun Zhang, Pukun Zhao, Fanqing Zhou, Longxiang Wang, and Haojian Huang. Memory-anchored multimodal reasoning for explainable video forensics. *arXiv preprint arXiv:2508.14581*, 2025a.
- Haoxing Chen, Yan Hong, Zizheng Huang, Zhuoer Xu, Zhangxuan Gu, Yaohui Li, Jun Lan, Huijia Zhu, Jianfu Zhang, Weiqiang Wang, and Huaxiong Li. Demamba: Ai-generated video detection on million-scale genvideo benchmark. *arXiv preprint arXiv:2405.19707*, 2024a.

- Jiaxin Chen, Miao Hu, Dengyong Zhang, and Jingyang Meng. Gc-consflow: Leveraging optical flow residuals and global context for robust deepfake detection. *arXiv preprint arXiv:2501.13435*, 2025b.
- Liang Chen, Yong Zhang, Yibing Song, Jue Wang, and Lingqiao Liu. Ost: Improving generalization of deepfake detection via one-shot test-time training. *Advances in Neural Information Processing Systems*, 35:24597–24610, 2022.
- Peng Chen, Jin Liu, Tao Liang, Guangzhi Zhou, Hongchao Gao, Jiao Dai, and Jizhong Han. Fsspotter: Spotting face-swapped video by spatial and temporal clues. In *2020 IEEE international conference on multimedia and expo (ICME)*, pages 1–6. IEEE Computer Society, 2020.
- Weiliang Chen, Wenzhao Zheng, Yu Zheng, Lei Chen, Jie Zhou, Jiwen Lu, and Yueqi Duan. Genworld: Towards detecting ai-generated real-world simulation videos, 2025c. <https://arxiv.org/abs/2506.10975>.
- Yize Chen, Zhiyuan Yan, Guangliang Cheng, Kangran Zhao, Siwei Lyu, and Baoyuan Wu. X2-dfd: A framework for explainable and extendable deepfake detection, 2024b. <https://arxiv.org/abs/2410.06126>.
- Yubin Chen, Xuyang Guo, Zhenmei Shi, Zhao Song, and Jiahao Zhang. T2vworldbench: A benchmark for evaluating world knowledge in text-to-video generation. *arXiv preprint arXiv:2507.18107*, 2025d. <https://arxiv.org/abs/2507.18107>.
- Zhikai Chen, Lingxi Xie, Shanmin Pang, Yong He, and Bo Zhang. Magdr: Mask-guided detection and reconstruction for defending deepfakes. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 9014–9023, June 2021. https://openaccess.thecvf.com/content/CVPR2021/html/Chen_MagDR_Mask-Guided_Detection_and_Reconstruction_for_Defending_Deepfakes_CVPR_2021_paper.html.
- Harry Cheng, Yangyang Guo, Tianyi Wang, Qi Li, Xiaojun Chang, and Liqiang Nie. Voice-face homogeneity tells deepfake, 2022a. <https://arxiv.org/abs/2203.02195>.
- Jikang Cheng, Zhiyuan Yan, Ying Zhang, Yuhao Luo, Zhongyuan Wang, and Chen Li. Can we leave deepfake data behind in training deepfake detector? *Advances in Neural Information Processing Systems*, 37:21979–21998, 2024.
- Kun Cheng, Xiaodong Cun, Yong Zhang, Menghan Xia, Fei Yin, Mingrui Zhu, Xuan Wang, Jue Wang, and Nannan Wang. Videoretalking: Audio-based lip synchronization for talking head video editing in the wild. In *SIGGRAPH Asia 2022 Conference Papers*, 2022b. <https://arxiv.org/abs/2211.14758>.
- Aditya Chinchure, Sahithya Ravi, Pushkar Shukla, Vered Shwartz, and Leonid Sigal. Spotlight: Identifying and localizing video generation errors using vlms. *arXiv preprint arXiv:2511.18102*, 2025. <https://arxiv.org/abs/2511.18102>.
- Jongwook Choi, Taehoon Kim, Yonghyun Jeong, Seungryul Baek, and Jongwon Choi. Exploiting style latent flows for generalizing deepfake video detection. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 1133–1143, 2024.

- Komal Chugh, Parul Gupta, Abhinav Dhall, and Ramanathan Subramanian. Not made for each other- audio-visual dissonance-based deepfake detection and localization. In Chang Wen Chen, Rita Cucchiara, Xian-Sheng Hua, Guo-Jun Qi, Elisa Ricci, Zhengyou Zhang, and Roger Zimmermann, editors, *MM 2020 - Proceedings of the 28th ACM International Conference on Multimedia*, MM 2020 - Proceedings of the 28th ACM International Conference on Multimedia, pages 439–447, United States, October 2020. Association for Computing Machinery (ACM). doi: 10.1145/3394171.3413700. Publisher Copyright: © 2020 ACM.; 28th ACM International Conference on Multimedia, MM 2020, ACM MM ; Conference date: 12-10-2020 Through 16-10-2020.
- Umur Aybars Ciftci, Ilke Demir, and Lijun Yin. Fakecatcher: Detection of synthetic portrait videos using biological signals. *IEEE transactions on pattern analysis and machine intelligence*, 2020a.
- Umur Aybars Ciftci, Ilke Demir, and Lijun Yin. How do the hearts of deep fakes beat? deep fake source detection via interpreting residuals with biological signals. In *2020 IEEE international joint conference on biometrics (IJCB)*, pages 1–10. IEEE, 2020b.
- Coalition for Content Provenance and Authenticity. C2pa technical specification, 2024. https://spec.c2pa.org/specifications/specifications/1.3/specs/C2PA_Specification.html. Accessed: 2026-01-01.
- Davide Alessandro Coccomini, Nicola Messina, Claudio Gennaro, and Fabrizio Falchi. Combining efficientnet and vision transformers for video deepfake detection. In *International conference on image analysis and processing*, pages 219–229. Springer, 2022.
- Riccardo Corvi, Davide Cozzolino, Ekta Prashnani, Shalini De Mello, Koki Nagano, and Luisa Verdoliva. Seeing what matters: Generalizable ai-generated video detection with forensic-oriented augmentation. *arXiv preprint arXiv:2506.16802*, 2025.
- Davide Cozzolino, Justus Thies, Gernot Riegler, Oliver Wang, Matthias Niessner, and Luisa Verdoliva. ID-Reveal: Identity-aware DeepFake Video Detection. In *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 15088–15097. IEEE, 2021. doi: 10.1109/ICCV48922.2021.01483. <https://doi.org/10.1109/ICCV48922.2021.01483>.
- Davide Cozzolino, Alessandro Pianese, Matthias Nießner, and Luisa Verdoliva. Audio-visual person-of-interest deepfake detection, 2023. <https://arxiv.org/abs/2204.03083>.
- Florinel-Alin Croitoru, Andrei-Iulian Hiji, Vlad Hondru, Nicolae Cătălin Ristea, Paul Irofti, Marius Popescu, Cristian Rusu, Radu Tudor Ionescu, Fahad Shahbaz Khan, and Mubarak Shah. Deepfake media generation and detection in the generative ai era: A survey and outlook. *arXiv preprint*, 2024. <https://arxiv.org/abs/2411.19537>.
- Florinel-Alin Croitoru, Vlad Hondru, Marius Popescu, Radu Tudor Ionescu, Fahad Shahbaz Khan, and Mubarak Shah. Mavos-dd: Multilingual audio-video open-set deepfake detection benchmark. *arXiv preprint arXiv:2505.11109*, 2025.
- Alessandro D’Amelio, Raffaella Lanzarotti, Sabrina Patania, Giuliano Grossi, Vittorio Cuculo, Andrea Valota, and Giuseppe Boccignone. On using rppg signals for deepfake detection: A cautionary note. In *Image Analysis and Processing – ICIAP 2023*, volume

- 14234 of *Lecture Notes in Computer Science*, pages 235–246. Springer, 2023. doi: 10.1007/978-3-031-43153-1_20.
- Soumyya Kanti Datta, Tanvi Ranga, Chengzhe Sun, and Siwei Lyu. Pia: Deepfake detection using phoneme-temporal and identity-dynamic analysis, 2025. <https://arxiv.org/abs/2510.14241>.
- Jingyi Deng, Chenhao Lin, Zhengyu Zhao, Shuai Liu, Zhe Peng, Qian Wang, and Chao Shen. A survey of defenses against ai-generated visual media: Detection, disruption, and authentication. *ACM Computing Surveys*, 58(5):1–35, 2025. doi: 10.1145/3770916. <http://dx.doi.org/10.1145/3770916>.
- Brian Dolhansky, Joanna Bitton, Ben Pflaum, Jikuo Lu, Russ Howes, Menglin Wang, and Cristian Canton Ferrer. The deepfake detection challenge (dfdc) dataset. *arXiv preprint arXiv:2006.07397*, 2020. <https://arxiv.org/abs/2006.07397>.
- Xiaoyi Dong, Jianmin Bao, Dongdong Chen, Weiming Zhang, Nenghai Yu, Dong Chen, Fang Wen, and Baining Guo. Identity-driven deepfake detection. *arXiv preprint arXiv:2012.03930*, 2020.
- Mengnan Du, Shiva Pentiyala, Yuening Li, and Xia Hu. Towards generalizable deepfake detection with locality-aware autoencoder. In *Proceedings of the 29th ACM international conference on information & knowledge management*, pages 325–334, 2020.
- Yuxuan Du, Zhendong Wang, Yuhao Luo, Caiyong Piao, Zhiyuan Yan, Hao Li, and Li Yuan. Cad: A general multimodal framework for video deepfake detection via cross-modal alignment and distillation, 2025. <https://arxiv.org/abs/2505.15233>.
- Chao Feng, Ziyang Chen, and Andrew Owens. Self-supervised video forensics by audio-visual anomaly detection. In *proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10491–10503, 2023.
- Steven Lawrence Fernandes, Raghavender Reddy O, Satoshi Oishi, Niussha Vosoughi, Subrahmanyam Mupparaju, and Utkarsh Mittal. Predicting heart rate variations of deepfake videos using neural ode. In *2019 IEEE/CVF International Conference on Computer Vision Workshops (ICCVW)*, pages 1721–1729. IEEE, 2019. doi: 10.1109/ICCVW.2019.00213. <https://doi.org/10.1109/ICCVW.2019.00213>.
- Xingyu Fu, Siyi Liu, Yinuo Xu, Pan Lu, Guangqiuse Hu, Tianbo Yang, Taran Anantasagar, Christopher Shen, Yikai Mao, Yuanzhe Liu, Keyush Shah, Chung Un Lee, Yejin Choi, James Zou, Dan Roth, and Chris Callison-Burch. Learning human-perceived fakeness in ai-generated videos via multimodal llms. *arXiv preprint arXiv:2509.22646*, 2025. <https://arxiv.org/abs/2509.22646>. Project page: <https://deeptracereward.github.io/>.
- Yifeng Gao, Yifan Ding, Hongyu Su, Juncheng Li, Yunhan Zhao, Lin Luo, Zixing Chen, Li Wang, Xin Wang, Yixu Wang, et al. David-xr1: Detecting ai-generated videos with explainable reasoning. *arXiv preprint arXiv:2506.14827*, 2025.
- Google DeepMind. Veo 3. <https://aistudio.google.com/models/veo-3>, 2025. Accessed: 2025-12-15.

- Jing Gu, Xin Wang, Hang Xu, Dong Wu, Yijin Chen, Xianghe Cong, Shaoyi Duan, Yuanfang Li, Chongyang Gao, Fukun Yin, Xun Li, Ziwei Liu, Jiaxing Huang, Yu Qiao, and Wei Gao. Phyworldbench: A physical realism benchmark for text-to-video generation. *arXiv preprint arXiv:2507.13428*, 2025. <https://arxiv.org/abs/2507.13428>.
- Zheyuan Gu, Qingsong Zhao, Yusong Wang, Zhaohong Huang, Xinqi Li, Cheng Yuan, Jiaowei Shao, Chi Zhang, and Xuelong Li. Beyond static artifacts: A forensic benchmark for video deepfake reasoning in vision language models. In *2026 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. doi: 10.48550/ARXIV.2602.21779. <https://doi.org/10.48550/ARXIV.2602.21779>. Submitted for publication.
- Zhihao Gu, Yang Chen, Taiping Yao, Shouhong Ding, Jilin Li, Feiyue Huang, and Lizhuang Ma. Spatiotemporal inconsistency learning for deepfake video detection. In *Proceedings of the 29th ACM international conference on multimedia*, pages 3473–3481, 2021.
- Zhihao Gu, Yang Chen, Taiping Yao, Shouhong Ding, Jilin Li, and Lizhuang Ma. Delving into the local: Dynamic inconsistency learning for deepfake video detection. In *Proceedings of the AAAI conference on artificial intelligence*, volume 36, pages 744–752, 2022a.
- Zhihao Gu, Taiping Yao, Yang Chen, Shouhong Ding, and Lizhuang Ma. Hierarchical contrastive inconsistency learning for deepfake video detection. In *European conference on computer vision*, pages 596–613. Springer, 2022b.
- Zhihao Gu, Taiping Yao, Yang Chen, Ran Yi, Shouhong Ding, and Lizhuang Ma. Region-aware temporal inconsistency learning for deepfake video detection. In *IJCAI*, volume 1, pages 920–926, 2022c.
- Xuyang Guo, Haoyu Zhou, Vivek Ramanujan, Chunhui Wang, Kang Li, Pengchuan Zhang, Kate Saenko, Kalyan Veeramachaneni, and X. Edward Zhou. T2vphysbench: A first-principles benchmark for physical consistency in text-to-video generation. *arXiv preprint arXiv:2505.00337*, 2025. <https://arxiv.org/abs/2505.00337>.
- Fahima Hajjej, Muhammad Hamid, and Ala Saleh Alluhaidan. An integrated framework for proactive deepfake mitigation via attention-driven watermarking and blockchain-based authenticity verification. *Scientific Reports*, 2026. doi: 10.1038/s41598-026-40166-6. <https://doi.org/10.1038/s41598-026-40166-6>. Early access article.
- Alexandros Haliassos, Konstantinos Vougioukas, Stavros Petridis, and Maja Pantic. Lips don’t lie: A generalisable and robust approach to face forgery detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5039–5049, 2021.
- Md. Tarek Hasan, Sanjay Saha, Shaojing Fan, Swakkhar Shatabda, and Terence Sim. Deepfake synthesis vs. detection: An uneven contest. *CoRR*, abs/2602.07986, 2026. doi: 10.48550/ARXIV.2602.07986. <https://doi.org/10.48550/ARXIV.2602.07986>.
- Ammarah Hashmi, Sahibzada Adil Shahzad, Chia-Wen Lin, Yu Tsao, and Hsin-Min Wang. Understanding audiovisual deepfake detection: Techniques, challenges, human factors and perceptual insights. *arXiv preprint*, 2024. <https://arxiv.org/abs/2411.07650>.

- Xinan He, Kaiqing Lin, Yue Zhou, Jiaming Zhong, Wei Ye, Wenhui Yi, Bing Fan, Feng Ding, Haodong Li, Bo Cao, and Bin Li. Mpf-net: Exposing high-fidelity ai-generated video forgeries via hierarchical manifold deviation and micro-temporal fluctuations. *CoRR*, abs/2601.21408, 2026. doi: 10.48550/ARXIV.2601.21408. <https://doi.org/10.48550/ARXIV.2601.21408>.
- Yinan He, Bei Gan, Siyu Chen, Yichun Zhou, Guojun Yin, Luchuan Song, Lu Sheng, Jing Shao, and Ziwei Liu. Forgerynet: A versatile benchmark for comprehensive forgery analysis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4360–4369, June 2021.
- Minji Heo and Simon S Woo. Fakechain: Exposing shallow cues in multi-step deepfake detection. In *Proceedings of the 34th ACM International Conference on Information and Knowledge Management*, pages 855–866, 2025.
- Javier Hernandez-Ortega, Ruben Tolosana, Julian Fierrez, and Aythami Morales. Deepfakeson-phys: Deepfakes detection based on heart rate estimation. *arXiv preprint arXiv:2010.00400*, 2020.
- Jonathan Ho, William Chan, Chitwan Saharia, Jay Whang, Ruiqi Gao, Alexey Gritsenko, Diederik P. Kingma, Ben Poole, Mohammad Norouzi, David J. Fleet, and Tim Salimans. Imagen video: High definition video generation with diffusion models, 2022a. <https://arxiv.org/abs/2210.02303>.
- Jonathan Ho, Tim Salimans, Alexey Gritsenko, William Chan, Mohammad Norouzi, and David J. Fleet. Video diffusion models, 2022b. <https://arxiv.org/abs/2204.03458>.
- Vlad Hondru, Eduard Hoge, Darian Onchis, and Radu Tudor Ionescu. Exddv: A new dataset for explainable deepfake detection in video. *arXiv preprint arXiv:2503.14421*, 2025.
- Fa-Ting Hong, Zunnan Xu, Zixiang Zhou, Jun Zhou, Xiu Li, Qin Lin, Qinglin Lu, and Dan Xu. Audio-visual controlled video diffusion with masked selective state spaces modeling for natural talking head generation, 2025. <https://arxiv.org/abs/2504.02542>.
- Xinming Hou, Mingming Yang, Wenxiang Jiao, Xing Wang, Zhaopeng Tu, and Wayne Xin Zhao. Coact: A global-local hierarchy for autonomous agent collaboration, 2024. <https://arxiv.org/abs/2406.13381>.
- Ziheng Hu, Hongtao Xie, Yuxin Wang, Jiahong Li, Zhongyuan Wang, and Yongdong Zhang. Dynamic inconsistency-aware deepfake video detection. In *IJCAI*, pages 736–742, 2021.
- Christian Interò, Robert Geirhos, Markus Olhofer, Sunny Liu, Barbara Hammer, and David Klindt. Ai-generated video detection via perceptual straightening. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. <https://openreview.net/forum?id=LsmUgStXby>.
- Lichuan Ji, Yingqi Lin, Zhenhua Huang, Yan Han, Xiaogang Xu, Jiafei Wu, Chong Wang, and Zhe Liu. Distinguish any fake videos: Unleashing the power of large-scale data and motion features. *arXiv preprint arXiv:2405.15343*, 2024.
- Shan Jia, Reilin Lyu, Kangran Zhao, Yize Chen, Zhiyuan Yan, Yan Ju, Chuanbo Hu, Xin Li,

- Baoyuan Wu, and Siwei Lyu. Can chatgpt detect deepfakes? a study of using multimodal large language models for media forensics. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 4324–4333, 2024. https://openaccess.thecvf.com/content/CVPR2024W/WMF/papers/Jia_Can_ChatGPT_Detect_DeepFakes_A_Study_of_Using_Multimodal_Large_CVPRW_2024_paper.pdf.
- Changjiang Jiang, Wenhui Dong, Zhonghao Zhang, Chenyang Si, Fengchang Yu, Wei Peng, Xinbin Yuan, Yifei Bi, Ming Zhao, Zian Zhou, et al. Ivy-fake: A unified explainable framework and benchmark for image and video aigc detection. *arXiv preprint arXiv:2506.00979*, 2025.
- Liming Jiang, Ren Li, Wayne Wu, Chen Qian, and Chen Change Loy. Deeperforensics-1.0: A large-scale dataset for real-world face forgery detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2020.
- Juho Jung, Sangyoun Lee, Jooeon Kang, and Yunjin Na. Www: Where, which and whatever enhancing interpretability in multimodal deepfake detection. *arXiv preprint arXiv:2408.02954*, 2024.
- Bingyi Kang, Yang Yue, Rui Lu, Zhijie Lin, Yang Zhao, Kaixin Wang, Gao Huang, and Jiashi Feng. How far is video generation from world model: A physical law perspective. *arXiv preprint arXiv:2411.02385*, 2024. <https://arxiv.org/abs/2411.02385>.
- Gautam Siddharth Kashyap, Harsh Joshi, Niharika Jain, Ebad Shabbir, Jiechao Gao, Nipun Joshi, and Usman Naseem. Revealing the truth with conllm for detecting multi-modal deepfakes. *CoRR*, abs/2601.17530, 2026. doi: 10.48550/ARXIV.2601.17530. <https://doi.org/10.48550/ARXIV.2601.17530>.
- Vinaya Sree Katamneni and Ajita Rattani. Contextual cross-modal attention for audio-visual deepfake detection and localization, 2024. <https://arxiv.org/abs/2408.01532>.
- Achhardeep Kaur, Azadeh Noori Hoshyar, Vidya Saikrishna, Selena Firmin, and Feng Xia. Deepfake video detection: challenges and opportunities. *Artificial Intelligence Review*, 57, 2024. doi: 10.1007/s10462-024-10810-6. <https://link.springer.com/article/10.1007/s10462-024-10810-6>.
- Hasam Khalid, Shahroz Tariq, Minha Kim, and Simon S. Woo. Fakeavceleb: A novel audio-video multimodal deepfake dataset. *arXiv preprint arXiv:2108.05080*, 2021. <https://arxiv.org/abs/2108.05080>.
- Taehoon Kim, Jongwook Choi, Yonghyun Jeong, Haeun Noh, Jaejun Yoo, Seungryul Baek, and Jongwon Choi. Beyond spatial frequency: Pixel-wise temporal frequency-based deepfake video detection. *arXiv preprint arXiv:2507.02398*, 2025.
- Youngseo Kim, Kwan Yun, Seokhyeon Hong, Sihun Cha, Colette Suhjung Koo, and Junyong Noh. X-AVDT: Audio-visual cross-attention for robust deepfake detection. In *2026 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. doi: 10.48550/ARXIV.2603.08483. <https://doi.org/10.48550/ARXIV.2603.08483>. Accepted as conference paper.

- Odin Kohler, Rahul Vijaykumar, and Masudul H Imtiaz. Deepfake detection in dyadic video calls using point of gaze tracking. *arXiv preprint arXiv:2509.25503*, 2025.
- Christos Koutlis and Symeon Papadopoulos. Dimodif: Discourse modality-information differentiation for audio-visual deepfake detection and localization, 2025. <https://arxiv.org/abs/2411.10193>.
- Kartik Kuckreja, Parul Gupta, Injy Hamed, Thamar Solorio, Muhammad Haris Khan, and Abhinav Dhall. Tell me habibi, is it real or fake? *arXiv preprint arXiv:2505.22581*, 2025.
- Ivan Kukanov and Jun Wah Ng. Klassify to verify: Audio-visual deepfake detection using ssl-based audio and handcrafted visual features, 2025. <https://arxiv.org/abs/2508.07337>.
- Rohit Kundu, Hao Xiong, Vishal Mohanty, Athula Balachandran, and Amit K Roy-Chowdhury. Towards a universal synthetic video detector: From face or background manipulations to fully ai-generated content. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 28050–28060, 2025.
- Patrick Kwon, Jaeseong You, Gyuhyeon Nam, Sungwoo Park, and Gyeongsu Chae. Kodf: A large-scale korean deepfake detection dataset. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 10744–10753, 2021.
- Ching-Yi Lai, Chiou ting Hsu, Chih-Chung Hsu, and Chia-Wen Lin. Prompt-guided multi-modal contrastive learning for cross-compression-rate deepfake detection. In *35th British Machine Vision Conference 2024, BMVC 2024, Glasgow, UK, November 25–28, 2024*. BMVA, 2024. <https://papers.bmvc2024.org/0619.pdf>.
- Nicolas Larue, Ngoc-Son Vu, Vitomir Struc, Peter Peer, and Vassilis Christophides. Seeable: Soft discrepancies and bounded contrastive learning for exposing deepfakes. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 21011–21021, 2023.
- Binh M Le and Simon S Woo. Quality-agnostic deepfake detection with intra-model collaborative learning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 22378–22389, 2023.
- Soo-Hyun Lee, Gyung-Eun Yun, Min Young Lim, and Youn Kyu Lee. A study on effective use of bpm information in deepfake detection. In *2021 International Conference on Information and Communication Technology Convergence (ICTC)*, pages 425–427. IEEE, 2021.
- Taegyeong Lee, Soyeong Kwon, and Taehwan Kim. Grid diffusion models for text-to-video generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8734–8743, 2024. <https://doi.org/10.1109/cvpr52733.2024.00834>.
- Roberto Leotta, Salvatore Alfio Sambataro, Claudio Vittorio Ragaglia, Mirko Casu, Yuri Petralia, Francesco Guarnera, Luca Guarnera, and Sebastiano Battiato. Synthforensics: A multi-generator benchmark for detecting synthetic video deepfakes. *CoRR*, abs/2602.04939, 2026. doi: 10.48550/ARXIV.2602.04939. <https://doi.org/10.48550/ARXIV.2602.04939>.

- Chuoqiao Li, Zhiwu Huang, Danda Pani Paudel, Yabin Wang, Mohamad Shahbazi, Xiaopeng Hong, and Luc Van Gool. A continual deepfake detection benchmark: Dataset, methods, and essentials. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 1339–1349, January 2023.
- Jieyu Li, Xin Zhang, and Joey Tianyi Zhou. Aegis: Authenticity evaluation benchmark for ai-generated video sequences. In *Proceedings of the 33rd ACM International Conference on Multimedia*, MM '25, page 13346–13353, New York, NY, USA, 2025a. Association for Computing Machinery. ISBN 9798400720352. doi: 10.1145/3746027.3758295. <https://doi.org/10.1145/3746027.3758295>.
- Qingcao Li, Miao He, Liang Yi, Qing Wen, Yitao Zhang, Hongshuo Jin, Peng Cheng, Zhongjie Ba, Li Lu, and Kui Ren. Divide and conquer: Multimodal video deepfake detection via cross-modal fusion and localization. *CoRR*, abs/2602.00209, 2026. doi: 10.48550/ARXIV.2602.00209. <https://doi.org/10.48550/ARXIV.2602.00209>.
- Tianqi Li, Ruobing Zheng, Minghui Yang, Jingdong Chen, and Ming Yang. Ditto: Motion-space diffusion for controllable realtime talking head synthesis, 2024a. <https://arxiv.org/abs/2411.19509>.
- Xiaodan Li, Yining Lang, Yuefeng Chen, Xiaofeng Mao, Yuan He, Shuhui Wang, Hui Xue, and Quan Lu. Sharp multiple instance learning for deepfake video detection. In *Proceedings of the 28th ACM international conference on multimedia*, pages 1864–1872, 2020a.
- Xiaolou Li, Zehua Liu, Chen Chen, Lantian Li, Li Guo, and Dong Wang. Zero-shot fake video detection by audio-visual consistency, 2024b. <https://arxiv.org/abs/2406.07854>.
- Yifei Li, Wenzhao Zheng, Yanran Zhang, Runze Sun, Yu Zheng, Lei Chen, Jie Zhou, and Jiwen Lu. Skyra: Ai-generated video detection via grounded artifact reasoning, 2025b. <https://arxiv.org/abs/2512.15693>.
- Yuezun Li, Ming-Ching Chang, and Siwei Lyu. In ictu oculi: Exposing ai created fake videos by detecting eye blinking. In *2018 IEEE International workshop on information forensics and security (WIFS)*, pages 1–7. Ieee, 2018.
- Yuezun Li, Xin Yang, Pu Sun, Honggang Qi, and Siwei Lyu. Celeb-df: A large-scale challenging dataset for deepfake forensics. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 3207–3216, 2020b.
- Zongxia Li, Xiyang Wu, Guangyao Shi, Yubin Qin, Hongyang Du, Tianyi Zhou, Dinesh Manocha, and Jordan Lee Boyd-Graber. Videohallu: Evaluating and mitigating multi-modal hallucinations on synthetic video understanding. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025c. <https://openreview.net/forum?id=NoC9HT7Kf7>.
- Yachao Liang, Min Yu, Gang Li, Jianguo Jiang, Boquan Li, Feng Yu, Ning Zhang, Xiang Meng, and Weiqing Huang. Speechforensics: Audio-visual speech representation learning for face forgery detection, 2025. <https://arxiv.org/abs/2508.09913>.
- Jianfeng Liao, Yichen Wei, Raymond Chan Ching Bon, Shulan Wang, Kam-Pui Chow, and Kwok-Yan Lam. Deepfake forensics adapter: A dual-stream network for generalizable

- deepfake detection. *CoRR*, abs/2603.01450, 2026. doi: 10.48550/ARXIV.2603.01450. <https://doi.org/10.48550/ARXIV.2603.01450>.
- Kaiqing Lin, Yuzhen Lin, Weixiang Li, Taiping Yao, and Bin Li. Standing on the shoulders of giants: Reprogramming visual-language model for general deepfake detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 5262–5270, 2025a. <https://ojs.aaai.org/index.php/AAAI/article/view/32559>.
- Li Lin, Neeraj Gupta, Yue Zhang, Hainan Ren, Chun-Hao Liu, Feng Ding, Xin Wang, Xin Li, Luisa Verdoliva, and Shu Hu. Detecting multimedia generated by large ai models: A survey. *arXiv preprint*, 2024a. doi: 10.48550/arXiv.2402.00045. <https://arxiv.org/abs/2402.00045>.
- Li Lin, Santosh Santosh, Mingyang Wu, Xin Wang, and Shu Hu. Ai-face: A million-scale demographically annotated ai-generated face dataset and fairness benchmark. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 3503–3515, 2025b.
- Yuzhen Lin, Wentang Song, Bin Li, Yuezun Li, Jiangqun Ni, Han Chen, and Qiushi Li. Fake it till you make it: Curricular dynamic forgery augmentations towards general deepfake detection. In *European Conference on Computer Vision*, pages 104–122, 2024b.
- Haoyu Liu, Chaoyu Gong, Mengke He, Jiate Li, Kai Han, and Siqiang Luo. When deepfake detection meets graph neural network: a unified and lightweight learning framework. *arXiv preprint arXiv:2508.05526*, 2025a.
- Jiaxin Liu, Jia Wang, Saihui Hou, Min Ren, Huijia Wu, Long Ma, Renwang Pei, and Zhaofeng He. Beyond face swapping: A diffusion-based digital human benchmark for multimodal deepfake detection. *arXiv preprint arXiv:2505.16512*, 2025b.
- Ping Liu, Qiqi Tao, and Joey Tianyi Zhou. Evolving from single-modal to multi-modal facial deepfake detection: Progress and challenges. *arXiv preprint*, 2025c. <https://arxiv.org/abs/2406.06965>.
- Qingyuan Liu, Pengyuan Shi, Yun-Yun Tsai, Chengzhi Mao, and Junfeng Yang. Turns out i’m not real: Towards robust detection of ai-generated videos. *arXiv preprint arXiv:2406.09601*, 2024a.
- Qingyuan Liu, Yun-Yun Tsai, Ruijian Zha, Victoria Li, Pengyuan Shi, Chengzhi Mao, and Junfeng Yang. Lavid: An agentic lvlm framework for diffusion-generated video detection. *arXiv preprint arXiv:2502.14994*, 2025d. <https://arxiv.org/abs/2502.14994>.
- X. Liu et al. Unleashing the potential of consistency learning for detecting and grounding multi-modal media manipulation. *arXiv preprint arXiv:2506.05890*, 2025e.
- Yuhan Liu, Xiuying Chen, Xiaoqing Zhang, Xing Gao, Ji Zhang, and Rui Yan. From skepticism to acceptance: simulating the attitude dynamics toward fake news. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence*, pages 7886–7894, 2024b.
- Yuhan Liu, Yuxuan Liu, Xiaoqing Zhang, Xiuying Chen, and Rui Yan. The truth becomes

- clearer through debate! multi-agent systems with large language models unmask fake news. In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 504–514, 2025f.
- Yuhan Liu, Zirui Song, Juntian Zhang, Xiaoqing Zhang, Xiuying Chen, and Rui Yan. The stepwise deception: Simulating the evolution from true news to fake news with llm agents. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 26187–26203, 2025g.
- Junxian Ma, Shiwen Wang, Jian Yang, Junyi Hu, Jian Liang, Guosheng Lin, Jingbo Chen, Kai Li, and Yu Meng. Sayanything: Audio-driven lip synchronization with conditional video diffusion, 2025. <https://arxiv.org/abs/2502.11515>.
- Long Ma, Zhiyuan Yan, Qinglang Guo, Yong Liao, Haiyang Yu, and Pengyuan Zhou. Detecting ai-generated video via frame consistency. *arXiv preprint arXiv:2402.02085*, 2024a.
- Long Ma, Jiajia Zhang, Hongping Deng, Ningyu Zhang, Yong Liao, and Haiyang Yu. Decof: Generated video detection via frame consistency. *CoRR*, 2024b.
- Long Ma, Zihao Xue, Yan Wang, Zhiyuan Yan, Jin Xu, Xiaorui Jiang, Haiyang Yu, Yong Liao, and Zhen Bi. Your one-stop solution for ai-generated video detection. *CoRR*, abs/2601.11035, 2026. doi: 10.48550/ARXIV.2601.11035. <https://doi.org/10.48550/ARXIV.2601.11035>.
- Maoyu Mao and Jun Yang. Exposing deepfake with pixel-wise ar and ppg correlation from faint signals. *arXiv preprint arXiv:2110.15561*, 2021.
- Iacopo Masi, Aditya Killekar, Royston Marian Mascarenhas, Shenoy Pratik Gurudatt, and Wael AbdAlmageed. Two-branch recurrent network for isolating deepfakes in videos. In *European conference on computer vision*, pages 667–684. Springer, 2020.
- Fanqing Meng, Jiaqi Liao, Xinyu Tan, Quanfeng Lu, Wenqi Shao, Kaipeng Zhang, Yu Cheng, Dianqi Li, Yu Qiao, and Ping Luo. Towards world simulator: Crafting physical commonsense-based benchmark for video generation. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025. <https://proceedings.mlr.press/v267/meng25c.html>.
- Hui Miao, Yuanfang Guo, Zeming Liu, and Yunhong Wang. Multi-modal deepfake detection via multi-task audio-visual prompt learning. In *Proceedings of the Thirty-Ninth AAI Conference on Artificial Intelligence and Thirty-Seventh Conference on Innovative Applications of Artificial Intelligence and Fifteenth Symposium on Educational Advances in Artificial Intelligence*, AAAI’25/IAAI’25/EAAI’25. AAAI Press, 2025. ISBN 978-1-57735-897-8. doi: 10.1609/aaai.v39i1.32042. <https://doi.org/10.1609/aaai.v39i1.32042>.
- Yisroel Mirsky and Wenke Lee. The creation and detection of deepfakes: A survey. *ACM Computing Surveys*, 54(1):1–41, 2021. doi: 10.1145/3425780. <http://dx.doi.org/10.1145/3425780>.
- Trisha Mittal, Uttaran Bhattacharya, Rohan Chandra, Aniket Bera, and Dinesh Manocha.

- Emotions don't lie: An audio-visual deepfake detection method using affective cues, 2020. <https://arxiv.org/abs/2003.06711>.
- Saman Motamed, Minghao Chen, Luc Van Gool, and Iro Laina. Travl: A recipe for making video-language models better judges of physics implausibility. *arXiv preprint arXiv:2510.07550*, 2025a. <https://arxiv.org/abs/2510.07550>.
- Saman Motamed, Laura Culp, Kevin Swersky, Priyank Jaini, and Robert Geirhos. Do generative video models understand physical principles? *arXiv preprint arXiv:2501.09038*, 2025b.
- Sneha Muppalla, Shan Jia, and Siwei Lyu. Integrating audio-visual features for multimodal deepfake detection, 2023. <https://arxiv.org/abs/2310.03827>.
- Kartik Narayan, Harsh Agarwal, Kartik Thakral, Surbhi Mittal, Mayank Vatsa, and Richa Singh. Df-platter: Multi-face heterogeneous deepfake dataset. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 9739–9748, June 2023.
- Dat Nguyen, Marcella Astrid, Anis Kacem, Enjie Ghorbel, and Djamila Aouada. Vulnerability-aware spatio-temporal learning for generalizable and interpretable deepfake video detection. *arXiv preprint arXiv:2501.01184*, 2025.
- Thanh Thi Nguyen, Quoc Viet Hung Nguyen, Dung Tien Nguyen, Duc Thanh Nguyen, Thien Huynh-The, Saeid Nahavandi, Thanh Tam Nguyen, Quoc-Viet Pham, and Cuong M. Nguyen. Deep learning for deepfakes creation and detection: A survey. *Computer Vision and Image Understanding*, 223:103525, 2022. doi: 10.1016/j.cviu.2022.103525. <http://dx.doi.org/10.1016/j.cviu.2022.103525>.
- Hong-Hanh Nguyen-Le, Van-Tuan Tran, Dinh-Thuc Nguyen, and Nhien-An Le-Khac. Passive deepfake detection across multi-modalities: A comprehensive survey. *arXiv preprint*, 2025. <https://arxiv.org/abs/2411.17911>.
- Zhenliang Ni, Qiangyu Yan, Mouxiao Huang, Tianning Yuan, Yehui Tang, Hailin Hu, Xinghao Chen, and Yunhe Wang. Genvidbench: A challenging benchmark for detecting ai-generated video. *arXiv preprint arXiv:2501.11340*, 2025.
- Fan Nie, Jiangqun Ni, Jian Zhang, Bin Zhang, and Weizhe Zhang. Dip: diffusion learning of inconsistency pattern for general deepfake detection. *IEEE Transactions on Multimedia*, 2024.
- Trevine Oorloff, Surya Koppiseti, Nicolò Bonettini, Divyaraj Solanki, Ben Colman, Yaser Yacoob, Ali Shahriyari, and Gaurav Bharaj. Avff: Audio-visual feature fusion for video deepfake detection, 2024. <https://arxiv.org/abs/2406.02951>.
- OpenAI. Sora: Creating video from text. <https://openai.com/index/sora/>, 2024a. Accessed: 2025-12-15.
- OpenAI. Sora 2. <https://sora.chatgpt.com>, 2024b. Accessed: 2025-12-15.
- Sudev Kumar Padhi, Harshit Kumar, Umesh Kashyap, and Sk Subidh Ali. De-fake: Style based anomaly deepfake detection. *arXiv preprint arXiv:2507.03334*, 2025.

- Kyoungjun Park, Yifan Yang, Juheon Yi, Shicheng Zheng, Yifei Shen, Dongqi Han, Caihua Shan, Muhammad Muaz, and Lili Qiu. Vidguard-r1: Ai-generated video detection and explanation via reasoning mllms and rl. *arXiv preprint arXiv:2510.02282*, 2025. <https://arxiv.org/abs/2510.02282>.
- William Peebles and Saining Xie. Scalable diffusion models with transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4172–4182, 2023. https://openaccess.thecvf.com/content/ICCV2023/papers/Peebles_Scalable_Diffusion_Models_with_Transformers_ICCV_2023_paper.pdf.
- Gan Pei, Jiangning Zhang, Menghan Hu, Zhenyu Zhang, Chengjie Wang, Yunsheng Wu, Guangtao Zhai, Jian Yang, Chunhua Shen, and Dacheng Tao. Deepfake generation and detection: A benchmark and survey. *arXiv preprint*, 2024. <https://arxiv.org/abs/2403.17881>.
- Wenshuo Peng, Gongxuan Wang, Tianmeng Yang, Chuanhao Li, Xiaojie Xu, Hui He, and Kaipeng Zhang. Svbench: Evaluation of video generation models on social reasoning. *arXiv preprint arXiv:2512.21507*, 2025. <https://arxiv.org/abs/2512.21507>.
- Hua Qi, Qing Guo, Felix Juefei-Xu, Xiaofei Xie, Lei Ma, Wei Feng, Yang Liu, and Jianjun Zhao. Deeprrhythm: Exposing deepfakes with attentional visual heartbeat rhythms. In *Proceedings of the 28th ACM international conference on multimedia*, pages 4318–4327, 2020.
- Yiran Qin, Zhelun Shi, Jiwen Yu, Xijun Wang, Enshen Zhou, Lijun Li, Zhenfei Yin, Xihui Liu, Lu Sheng, Jing Shao, Lei Bai, Wanli Ouyang, and Ruimao Zhang. Worldsimbench: Towards video generation models as world simulators. In *Forty-Second International Conference on Machine Learning*, 2025. https://openreview.net/forum?id=4_hExrLJ7j.
- Andreas Rossler, Davide Cozzolino, Luisa Verdoliva, Christian Riess, Justus Thies, and Matthias Nießner. Faceforensics++: Learning to detect manipulated facial images. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 1–11, 2019.
- A. S. M. Sharifuzzaman Sagar, Mohammed Bennamoun, Farid Boussaid, Naeha Sharif, Lian Xu, Shaaban Sahmoud, and Ali Kishk. Fact or fake? assessing the role of deepfake detectors in multimodal misinformation detection. *CoRR*, abs/2602.01854, 2026. doi: 10.48550/ARXIV.2602.01854. <https://doi.org/10.48550/ARXIV.2602.01854>.
- Adrian Serrano, Erwan Umlil, and Ronan Thomas. Deepfake detectors are DUMB: A benchmark to assess adversarial training robustness under transferability constraints. *CoRR*, abs/2601.05986, 2026. doi: 10.48550/ARXIV.2601.05986. <https://doi.org/10.48550/ARXIV.2601.05986>.
- Sahibzada Adil Shahzad, Ammarah Hashmi, Sarwar Khan, Yan-Tsung Peng, Yu Tsao, and Hsin-Min Wang. Lip sync matters: A novel multimodal forgery detector. In *2022 Asia-Pacific Signal and Information Processing Association Annual Summit and Conference (APSIPA ASC)*, pages 1885–1892, 2022. doi: 10.23919/APSIPAASC55919.2022.9980296.
- Sahibzada Adil Shahzad, Ammarah Hashmi, Yan-Tsung Peng, Yu Tsao, and Hsin-Min Wang. How good is chatgpt at audiovisual deepfake detection: A comparative study of chatgpt, ai models and human perception. *APSIPA Transactions on Signal and Information*

- Processing*, 14:e11, 2025. doi: 10.1561/116.20250004. <https://www.nowpublishers.com/article/OpenAccessDownload/SIP-20250004>.
- Rui Shao, Tianxing Wu, Liqiang Nie, and Ziwei Liu. Deepfake-adapter: Dual-level adapter for deepfake detection. *International Journal of Computer Vision*, 133(6):3613–3628, January 2025. ISSN 1573-1405. doi: 10.1007/s11263-024-02274-6. <http://dx.doi.org/10.1007/s11263-024-02274-6>.
- Guangyu Shen, Zhihua Li, Xiang Xu, Tianchen Zhao, Zheng Zhang, Dongsheng An, Zhuowen Tu, Yifan Xing, and Qin Zhang. Authguard: Generalizable deepfake detection via language guidance. *arXiv preprint arXiv:2506.04501*, 2025. <https://arxiv.org/abs/2506.04501>.
- Danqing Shi, Lan Jiang, Katherine M. Collins, Shangzhe Wu, Ayush Tewari, and Miri Zilka. How do people watch AI-generated videos of physical scenes? *CoRR*, abs/2602.03374, 2026. doi: 10.48550/ARXIV.2602.03374. <https://doi.org/10.48550/ARXIV.2602.03374>.
- Uriel Singer, Adam Polyak, Thomas Hayes, Xi Yin, Jie An, Songyang Zhang, Qiyuan Hu, Harry Yang, Oron Ashual, Oran Gafni, Devi Parikh, Sonal Gupta, and Yaniv Taigman. Make-a-video: Text-to-video generation without text-video data, 2022. <https://arxiv.org/abs/2209.14792>.
- Stefan Smeu, Dragos-Alexandru Boldisor, Dan Oneata, and Elisabeta Oneata. Circumventing shortcuts in audio-visual deepfake detection datasets with unsupervised learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18815–18825, June 2025.
- Wentang Song, Zhiyuan Yan, Yuzhen Lin, Taiping Yao, Changsheng Chen, Shen Chen, Yandan Zhao, Shouhong Ding, and Bin Li. A quality-centric framework for generic deepfake detection. *arXiv preprint arXiv:2411.05335*, 2024a.
- Xiufeng Song, Xiao Guo, Jiache Zhang, Qirui Li, Lei Bai, Xiaoming Liu, Guangtao Zhai, and Xiaohong Liu. On learning multi-modal forgery representation for diffusion generated video detection. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 122054–122077. Curran Associates, Inc., 2024b. doi: 10.52202/079017-3878. https://proceedings.neurips.cc/paper_files/paper/2024/file/dccbeb7a8df3065c4646928985edf435-Paper-Conference.pdf.
- Sotirios Stamnas and Victor Sanchez. Difffake: Exposing deepfakes using differential anomaly detection. In *Proceedings of the Winter Conference on Applications of Computer Vision*, pages 695–705, 2025.
- Kalin Stefanov, Bhawna Paliwal, and Abhinav Dhall. Visual representations of physiological signals for fake video detection. *arXiv preprint arXiv:2207.08380*, 2022.
- Haoran Sun, Chen Cai, Huiping Zhuang, Kong Aik Lee, Lap-Pui Chau, and Yi Wang. Edvd-llama: Explainable deepfake video detection via multimodal large language model reasoning. *arXiv preprint arXiv:2510.16442*, 2025.

- Zekun Sun, Yujie Han, Zeyu Hua, Na Ruan, and Weijia Jia. Improving the efficiency and robustness of deepfakes detection through precise geometric features. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3609–3618, 2021.
- Hao Tan, Jun Lan, Zichang Tan, Ajian Liu, Chuanbiao Song, Senyuan Shi, Huijia Zhu, Weiqiang Wang, Jun Wan, and Zhen Lei. Veritas: Generalizable deepfake detection via pattern-aware reasoning. *arXiv preprint arXiv:2508.21048*, 2025. <https://arxiv.org/abs/2508.21048>.
- Hao Tan, Jun Lan, Senyuan Shi, Zichang Tan, Zijian Yu, Huijia Zhu, Weiqiang Wang, Jun Wan, and Zhen Lei. Videoveritas: AI-generated video detection via perception pretext reinforcement learning. *CoRR*, abs/2602.08828, 2026. doi: 10.48550/ARXIV.2602.08828. <https://doi.org/10.48550/ARXIV.2602.08828>.
- Shahroz Tariq, Sangyup Lee, and Simon Woo. One detector to rule them all: Towards a general deepfake attack detection framework. In *Proceedings of the web conference 2021*, pages 3625–3637, 2021.
- Jiahe Tian, Cai Yu, Xi Wang, Peng Chen, Zihao Xiao, Jiao Dai, Jizhong Han, and Yesheng Chai. Real appearance modeling for more general deepfake detection. In *European Conference on Computer Vision*, pages 402–419. Springer, 2024.
- Ruben Tolosana, Ruben Vera-Rodriguez, Julian Fierrez, Aythami Morales, and Javier Ortega-Garcia. Deepfakes and beyond: A survey of face manipulation and fake detection. *Information Fusion*, 64:131–148, 2020. doi: 10.1016/j.inffus.2020.06.014. <http://dx.doi.org/10.1016/j.inffus.2020.06.014>.
- Timothy Tong, David Anastasiu, and Yuhong Liu. Deepfake detection using spatiotemporal methods and vision-language models. *Proceedings of the KDD Undergraduate and Master’s Consortium (KDD-UMC’25)*, 3, 2025.
- Loc Trinh, Michael Tsang, Sirisha Rambhatla, and Yan Liu. Interpretable and trustworthy deepfake detection via dynamic prototypes. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pages 1973–1983, 2021.
- Eleanor Tursman, Marilyn George, Seny Kamara, and James Tompkin. Towards untrusted social video verification to combat deepfakes via face geometry consistency. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition workshops*, pages 654–655, 2020.
- Danial Samadi Vahdati, Tai D Nguyen, Aref Azizpour, and Matthew C Stamm. Beyond deepfake images: Detecting ai-generated videos. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4397–4408, 2024.
- Keerthi Veeramachaneni, Praveen Tirupattur, Amrit Singh Bedi, and Mubarak Shah. Leveraging pre-trained visual models for ai-generated video detection. *arXiv preprint arXiv:2507.13224*, 2025.
- Luisa Verdoliva. Media forensics and deepfakes: An overview. *IEEE Journal of Selected*

- Topics in Signal Processing*, 14(5):910–932, 2020. doi: 10.1109/JSTSP.2020.3002101. <http://dx.doi.org/10.1109/JSTSP.2020.3002101>.
- Ganglai Wang, Peng Zhang, Lei Xie, Wei Huang, Yufei Zha, and Yanning Zhang. An audio-visual attention based multimodal network for fake talking face videos detection, 2022a. <https://arxiv.org/abs/2203.05178>.
- H. Wang et al. T³svfnd: Towards an evolving fake news detector for emergencies with test-time training on short video platforms. *arXiv preprint arXiv:2507.20286*, 2025a.
- Hanyi Wang, Zihan Liu, and Shilin Wang. Exploiting complementary dynamic incoherence for deepfake video detection. *IEEE Transactions on Circuits and Systems for Video Technology*, 33(8):4027–4040, 2023.
- Jian Wang, Baoyuan Wu, Li Liu, and Qingshan Liu. Fauforensics: Boosting audio-visual deepfake detection with facial action units, 2025b. <https://arxiv.org/abs/2505.08294>.
- Jiaqi Wang, Weijia Wu, Yi Zhan, Rui Zhao, Ming Hu, James Cheng, Wei Liu, Philip Torr, and Kevin Qinghong Lin. Video reality test: Can ai-generated asmr videos fool vlms and humans?, 2025c. <https://arxiv.org/abs/2512.13281>.
- Runqi Wang, Yang Chen, Sijie Xu, Tianyao He, Wei Zhu, Dejie Song, Nemo Chen, Xu Tang, and Yao Hu. Dynamicface: High-quality and consistent face swapping for image and video using composable 3d facial priors. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 13438–13447, 2025d.
- Tianyi Wang and Kam Pui Chow. Noise based deepfake detection via multi-head relative-interaction. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 14548–14556, 2023.
- Tianyi Wang, Xin Liao, Kam Pui Chow, Xiaodong Lin, and Yinglong Wang. Deepfake detection: A comprehensive survey from the reliability perspective. *ACM Computing Surveys*, 2024a. doi: 10.1145/3699710. <https://dl.acm.org/doi/10.1145/3699710>.
- Xueyu Wang, Jiajun Huang, Siqi Ma, Surya Nepal, and Chang Xu. Deepfake disrupter: The detector of deepfake is my friend. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14920–14929, 2022b.
- Yaohua Wang, Siying Cui, Aixi Zhang, Wei-Long Zheng, Senzhang Wang, et al. Fuseany-part: Diffusion-driven facial parts swapping via multiple reference images. *Advances in Neural Information Processing Systems*, 37:80864–80884, 2024b.
- Yiping Wang, Xuehai He, Kuan Wang, Luyao Ma, Jianwei Yang, Shuohang Wang, Simon Shaolei Du, and Yelong Shen. Is your world simulator a good story presenter? a consecutive events-based benchmark for future long video generation. *arXiv preprint arXiv:2412.16211*, 2024c. <https://arxiv.org/abs/2412.16211>.
- Yubo Wang, Juntian Zhang, Yichen Wu, Yankai Lin, Nils Lukas, and Yuhan Liu. For-est before trees: Latent superposition for efficient visual reasoning. *arXiv preprint arXiv:2601.06803*, 2026.

- Yuxi Wang, Yikang Wang, Qishan Zhang, Hiromitsu Nishizaki, and Ming Li. Vcapav: A video-caption based audio-visual deepfake detection dataset. In *Proc. Interspeech 2025*, pages 3908–3912, 2025e.
- Zeqing Wang, Xinyu Wei, Bairui Li, Zhen Guo, Jinrui Zhang, Hongyang Wei, Keze Wang, and Lei Zhang. Videoverse: How far is your t2v generator from a world model? *arXiv preprint arXiv:2510.08398*, 2025f. <https://arxiv.org/abs/2510.08398>.
- Zeqing Wang, Xiaohui Yuan, Shiwei Ding, Debing Ye, Zizheng Chen, and Lin Chen. Phynetex: A benchmark dataset and method for detecting and explaining physical plausibility in text-to-video models. *arXiv preprint arXiv:2512.01843*, 2025g. <https://arxiv.org/abs/2512.01843>.
- Haiquan Wen, Yiwei He, Zhenglin Huang, Tianxiao Li, Zihan Yu, Xingru Huang, Lu Qi, Baoyuan Wu, Xiangtai Li, and Guangliang Cheng. Busterx: Mllm-powered ai-generated video forgery detection and explanation. *arXiv preprint arXiv:2505.12620*, 2025a.
- Haiquan Wen, Tianxiao Li, Zhenglin Huang, Yiwei He, and Guangliang Cheng. Busterx++: Towards unified cross-modal ai-generated content detection and explanation with mllm. *arXiv preprint arXiv:2507.14632*, 2025b. <https://arxiv.org/abs/2507.14632>.
- Jiahui Wu, Yu Zhu, Xiaoben Jiang, Yatong Liu, and Jiajun Lin. Local attention and long-distance interaction of rppg for deepfake detection. *The Visual Computer*, 40(2):1083–1094, 2024.
- Xuecheng Wu, Danlei Huang, Heli Sun, Xinyi Yin, Yifan Wang, Hao Wang, Jia Zhang, Fei Wang, Peihao Guo, Suyu Xing, Junxiao Xue, and Liang He. Hola: Enhancing audio-visual deepfake detection via hierarchical contextual aggregations and efficient pre-training, 2025. <https://arxiv.org/abs/2507.22781>.
- Qiang Xu, Wenpeng Mu, Jianing Li, Tanfeng Sun, and Xinghao Jiang. Advancements in ai-generated content forensics: A systematic literature review. *ACM Computing Surveys*, 58(3):1–36, 2025a. doi: 10.1145/3760526. <http://dx.doi.org/10.1145/3760526>.
- Wenbo Xu, Wei Lu, and Xiangyang Luo. Weakly supervised multimodal temporal forgery localization via multitask learning, 2025b. <https://arxiv.org/abs/2508.02179>.
- Wenbo Xu, Junyan Wu, Wei Lu, Xiangyang Luo, and Qian Wang. A multimodal deviation perceiving framework for weakly-supervised temporal forgery localization. In *Proceedings of the 33rd ACM International Conference on Multimedia*, MM '25, page 11581–11589. ACM, October 2025c. doi: 10.1145/3746027.3755534. <http://dx.doi.org/10.1145/3746027.3755534>.
- Ying Xu, Marius Pedersen, and Kiran Raja. Vod: Learning volume of differences for video-based deepfake detection. *arXiv preprint arXiv:2503.07607*, 2025d.
- Yuting Xu, Jian Liang, Gengyun Jia, Ziming Yang, Yanhao Zhang, and Ran He. Tall: Thumbnail layout for deepfake video detection. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 22658–22668, 2023.

- Yuting Xu, Jian Liang, Lijun Sheng, and Xiao-Yu Zhang. Learning spatiotemporal inconsistency via thumbnail layout for face deepfake detection. *International Journal of Computer Vision*, 132(12):5663–5680, 2024.
- Zhiyuan Yan, Yong Zhang, Xinhang Yuan, Siwei Lyu, and Baoyuan Wu. Deepfakebench: A comprehensive benchmark of deepfake detection. In A. Oh, T. Neumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 4534–4565. Curran Associates, Inc., 2023. https://proceedings.neurips.cc/paper_files/paper/2023/file/0e735e4b4f07de483cbe250130992726-Paper-Datasets_and_Benchmarks.pdf.
- Zhiyuan Yan, Tao Liu, Mei-Lin Zhang, Wenbin Li, Jianmin Zhang, and Jun Lu. Transcending forgery specificity with latent space augmentation for generalizable deepfake detection. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 8984–8994, 2024. doi: 10.1109/CVPR52733.2024.00858. <https://doi.org/10.1109/CVPR52733.2024.00858>.
- Zhiyuan Yan, Yandan Zhao, Shen Chen, Mingyi Guo, Xinghe Fu, Taiping Yao, Shouhong Ding, Yunsheng Wu, and Li Yuan. Generalizing deepfake video detection with plug-and-play: Video-level blending and spatiotemporal adapter tuning. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 12615–12625, 2025.
- Chenzhao Yang, Jun Ma, Shilin Wang, and Alan Wee-Chung Liew. Preventing deepfake attacks on speaker authentication by dynamic lip movement analysis. *IEEE Transactions on Information Forensics and Security*, 16:1841–1854, 2021. <https://api.semanticscholar.org/CorpusID:230998982>.
- Qingyang Yang, Chuanxu Wang, Peng Liu, Zitai Jiang, and Jiajiong Li. Video anomaly detection via self-supervised and spatio-temporal proxy tasks learning. *Pattern Recognition*, 158:111021, 2025.
- Xin Yang, Yuezun Li, and Siwei Lyu. Exposing deep fakes using inconsistent head poses. In *ICASSP 2019-2019 IEEE international conference on acoustics, speech and signal processing (ICASSP)*, pages 8261–8265. IEEE, 2019.
- Junyan Ye, Baichuan Zhou, Zilong Huang, Junan Zhang, Tianyi Bai, Hengrui Kang, Jun He, Honglin Lin, Zihao Wang, Tong Wu, et al. Loki: A comprehensive synthetic data detection benchmark using large multimodal models. *ICLR*, 2025.
- Qilin Yin, Wei Lu, Bin Li, and Jiwu Huang. Dynamic difference learning with spatio-temporal correlation for deepfake video detection. *IEEE Transactions on Information Forensics and Security*, 18:4046–4058, 2023.
- Peipeng Yu, Jianwei Fei, Hui Gao, Xuan Feng, Zhihua Xia, and Chip Hong Chang. Unlocking the capabilities of large vision-language models for generalizable and explainable deepfake detection. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 72925–72943. PMLR, 13–19 Jul 2025. <https://proceedings.mlr.press/v267/yu25d.html>.
- Sayeem Been Zaman, Wasimul Karim, Arefin Ittesafun Abian, Reem E. Mohamed,

- Md Rafiqul Islam, Asif Karim, and Sami Azam. Deepagent: A dual stream multi agent fusion for robust multimodal deepfake detection, 2025. <https://arxiv.org/abs/2512.07351>.
- Chenyu Zhang, Xunyu Wang, Seunghoon Hur, Yung-Hsu Yang, Shih-Lun Wu, Fanjie Wu, and Junsong Yuan. Morpheus: Benchmarking physical reasoning of video generative models with real physical experiments. *arXiv preprint arXiv:2504.02918*, 2025a. <https://arxiv.org/abs/2504.02918>.
- Daichi Zhang, Chenyu Li, Fanzhao Lin, Dan Zeng, and Shiming Ge. Detecting deepfake videos with temporal dropout 3dcnn. In *IJCAI*, pages 1288–1294, 2021a.
- Daichi Zhang, Zihao Xiao, Shikun Li, Fanzhao Lin, Jianmin Li, and Shiming Ge. Learning natural consistency representation for face forgery video detection. In *European Conference on Computer Vision*, pages 407–424. Springer, 2024a.
- David Junhao Zhang, Jay Zhangjie Wu, Jia-Wei Liu, Rui Zhao, Lingmin Ran, Yuchao Gu, Difei Gao, and Mike Zheng Shou. Show-1: Marrying pixel and latent diffusion models for text-to-video generation. *International Journal of Computer Vision*, 133(4):1879–1893, 2025b. <https://doi.org/10.1007/s11263-024-02271-9>.
- Jian Zhang, Jiangqun Ni, and Hao Xie. Deepfake videos detection using self-supervised decoupling network. In *2021 IEEE International Conference on Multimedia and Expo (ICME)*, pages 1–6, 2021b. doi: 10.1109/ICME51207.2021.9428368.
- Juntian Zhang, Chuanqi Cheng, Yuhan Liu, Wei Liu, Jian Luan, and Rui Yan. Weaving context across images: Improving vision-language models through focus-centric visual chains. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 27782–27798, 2025c.
- Juntian Zhang, Song Jin, Chuanqi Cheng, Yuhan Liu, Yankai Lin, Xun Zhang, Yufei Zhang, Fei Jiang, Guojun Yin, Wei Lin, et al. Viper: Empowering the self-evolution of visual perception abilities in vision-language model. *arXiv preprint arXiv:2510.24285*, 2025d.
- Qin Zhang, Peiyu Jing, Hong-Xing Yu, Fangqiang Ding, Fan Nie, Weimin Wang, Yilun Du, James Zou, Jiajun Wu, and Bing Shuai. Physion-eval: Evaluating physical realism in generated video via human reasoning. *arXiv preprint arXiv:2603.19607*, 2026. <https://arxiv.org/abs/2603.19607>.
- Shuhai Zhang, ZiHao Lian, Jiahao Yang, Daiyuan Li, Guoxuan Pang, Feng Liu, Bo Han, Shutao Li, and Mingkui Tan. Physics-driven spatiotemporal modeling for ai-generated video detection. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025e. <https://openreview.net/forum?id=HiBoJLCyEo>.
- Y. Zhang et al. Consistency-aware fake videos detection on short video platforms. *arXiv preprint arXiv:2504.21495*, 2025f.
- Yue Zhang, Ben Colman, Xiao Guo, Ali Shahriyari, and Gaurav Bharaj. Common sense reasoning for deepfake detection. In *European conference on computer vision*, pages 399–415. Springer, 2024b.

- Chende Zheng, Ruiqi Suo, Chenhao Lin, Zhengyu Zhao, Le Yang, Shuai Liu, Minghui Yang, Cong Wang, and Chao Shen. D3: Training-free ai-generated video detection using second-order features. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 12852–12862, 2025.
- Jiaran Zhou, Yuezun Li, Baoyuan Wu, Bin Li, Junyu Dong, et al. Freqblender: Enhancing deepfake detection by blending frequency knowledge. *Advances in Neural Information Processing Systems*, 37:44965–44988, 2024.
- Yipin Zhou and Ser-Nam Lim. Joint audio-visual deepfake detection. In *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 14780–14789, 2021. doi: 10.1109/ICCV48922.2021.01453.
- Bojia Zi, Minghao Chang, Jingjing Chen, Xingjun Ma, and Yu-Gang Jiang. Wilddeepfake: A challenging real-world dataset for deepfake detection. In *Proceedings of the 28th ACM International Conference on Multimedia*, MM '20, page 2382–2390, New York, NY, USA, 2020. Association for Computing Machinery. ISBN 9781450379885. doi: 10.1145/3394171.3413769. <https://doi.org/10.1145/3394171.3413769>.
- Heqing Zou, Meng Shen, Yuchen Hu, Chen Chen, Eng Siong Chng, and Deepu Rajan. Cross-modality and within-modality regularization for audio-visual deepfake detection, 2024. <https://arxiv.org/abs/2401.05746>.
- Yueying Zou, Peipei Li, Zekun Li, Huaibo Huang, Xing Cui, Xuannan Liu, Chenghanyu Zhang, and Ran He. Survey on ai-generated media detection: From non-mlm to mllm. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2502.05240. <https://arxiv.org/abs/2502.05240>.